

PROPOSED INTEGRATED DRAFT — NOT APPROVED*For Council review only · Not authorised, effective or operational until formally approved*

AIG-AGT-02 · The Council · AI GOVERNANCE

Agentic Classification Reference

Apply the classification profile to the specific configuration and UC-ID/outcome being assessed, not only to the shared system. Distinct purposes, populations, action scopes or permissions may have distinct risk and delegated decisions under the same AIR-ID. UC-ID applies to ordinary uses too; action-capability classifications apply only when that configuration can take or initiate actions. A profile or map is not approval.

Capability vector, autonomy scale and agency tiers

Field	Controlled value
Document title	Agentic Classification Reference
Status	Working draft; pre-adoption
Author	AI Governance Lead
Date	19 September 2026
Purpose	The shared vocabulary the Agent Record and the Agentic Triage both use, so a level or a capability means the same thing everywhere
Used by	Agent Record / ASBOM (Capability Vector and fields); Agentic Triage (agency profile and tiering)
Source	Uploaded agentic analysis; aligned to OWASP 2026 guidance (ASI03, ASI02, ASI08; Agent Control Standard)
Classification	OFFICIAL

1. Why one reference

The Agent Record and the Agentic Triage both classify agents, and they must speak the same language. This reference fixes the definitions once. It has no logic of its own; it is the dictionary the other artefacts point to.

The organising idea, from which everything else follows: cognitive capability is not agency. An agent can be brilliant at reasoning and pose little danger, while a mediocre model wired to a production database, payments and email can pose a great deal. So we separate what an agent can think from what it can do, and we govern the doing.

2. Cognitive capability versus action capability

Keep the two apart when describing any agent.

- Cognitive (thinking): reason, plan, analyse, generate, retrieve, infer, recommend, decide. These raise quality and error questions.
- Action (doing): read, write, execute, communicate, purchase, modify, delegate, deploy. These raise agency and consequence questions, and are where governance effort belongs.

3. The capability vector

Record what the agent can technically do, capability by capability, as No, Yes or Scoped. Capability is not intention: an instruction not to do something is not a control if the ability remains.

Capability	Governance question
Read	What information can it access?

Capability	Governance question
Write	What can it change?
Execute	What can it cause to happen?
Communicate	Who can it contact?
Purchase	Can it spend money?
Delegate	Can it create or subcontract another agent?
Persuade	Can it influence humans?
Code	Can it generate or execute software?
Discover	Can it independently find new resources?
Persist	Can it operate beyond a session?
Replicate	Can it create additional instances?
Learn	Can its behaviour change from experience?
Escalate	Can it request greater authority?

4. Agency multipliers

These capabilities change the risk picture sharply and are captured as explicit flags on the Agent Record, because any one of them can turn a modest agent into a high-agency one.

- The ten flags: persistence, memory, delegation, tool discovery, credential access, self-modification, replication, goal adaptation, external communication and financial authority.

5. The autonomy scale (A0 to A5)

Autonomy is permissioned, granted deliberately at a gate, not inferred from how sophisticated the system looks. The level is recorded on the Agent Record and revisited at reauthorisation.

Level	Meaning
A0 advisory	Agent recommends; a human acts.
A1 assisted	Agent prepares actions; a human explicitly executes.
A2 bounded execution	Agent executes pre-approved actions within strict limits.
A3 conditional autonomy	Agent selects actions within a defined policy envelope.
A4 delegated autonomy	Agent plans, selects tools, executes and adapts with limited human intervention.
A5 open-ended autonomy	Agent independently pursues complex objectives across changing environments.

6. The agency tiers (T0 to T5)

The tier is the output of the Agentic Triage's agency profile, and it routes the agent to a governance pathway. It is based on actual behaviour and authority, not the marketing description.

Tier	Meaning
T0 informational	No meaningful external agency.
T1 assisted	A human performs the consequential action.
T2 bounded agent	Acts, but within tightly constrained authority.

Tier	Meaning
T3 consequential agent	Can independently produce material external consequences.
T4 high-agency	Broad authority plus persistence plus significant autonomy or delegation.
T5 exceptional / high-consequence	High autonomy, broad capability, severe consequences or difficult containment.

7. The five agency dimensions

The triage scores these five and keeps them as a profile rather than adding them up, so a spike on one dimension is never hidden by low scores on the others. Controllability runs the opposite way, because low controllability is the danger.

Dimension	Scale (0 to 5)
Consequence	0 trivial to 5 potentially catastrophic or irreversible. What happens if it gets it wrong?
Autonomy	0 no action to 5 open-ended or persistent autonomy. How independently can it act?
Authority	0 information only to 5 highly consequential authority. How much real-world power does it hold?
Reach	0 isolated to 5 broad, public or physical. How much of the environment can it affect?
Controllability	Reversed: 0 instant stop with full rollback to 5 effectively irreversible. How easily can humans stop or reverse it?

8. How they combine

Two rules govern how these pieces are used together.

- Profile, not sum: the five dimensions are read as a profile, not a sum. A system with modest autonomy but high authority and low controllability deserves attention a single averaged score would bury.
- Intensity: governance intensity scales with agency times capability times consequence times uncertainty. This is why the same agent can be low risk summarising documents and high risk moving money: the action and authority changed, not the agent.

9. Status and source

Working draft, pre-adoption. The capability vector, the multipliers, the A0 to A5 scale, the T0 to T5 tiers and the five dimensions are the framing of the uploaded agentic analysis, adopted here as the common vocabulary and aligned to OWASP 2026 guidance (identity and privilege ASI03, tool misuse ASI02, cascading failures ASI08, and the Agent Control Standard), verified against genai.owasp.org on 19 September 2026.

Working draft. The dictionary the Agent Record and the Agentic Triage both point to.
Deterministic classification and escalation rules

Agency tier is a profile, not an average. The highest applicable mandatory floor governs. The governance decision record must document the profile and any escalation or exception.

- T0–T1: informational/assisted activity with no independent consequential action.
- T2: bounded agent with explicit scoped execution and reliable containment.
- T3: consequential agent where actions can materially affect people, records, services, finances or external communications.

- T4: high-agency operation involving significant persistence, delegation, broad tool access, high consequence or difficult containment.
- T5: exceptional/high-consequence agency, including unbounded or difficult-to-contain autonomy, agent creation/replication, self-modification, physical/OT reach or similarly exceptional combinations. Mandatory escalation triggers include: authority ceiling cannot be reliably enforced; containment cannot be demonstrated; delegation/recursion is unbounded; irreversible high-consequence actions are possible without appropriate human control; persistent/shared memory lacks effective isolation; or the agent can create or modify agents/configurations without separately authorised controls. Frontier model is recorded separately from agency tier and never substitutes for consequence assessment.

Annex — Classification identity and routing boundaries

PROPOSED INTEGRATION ONLY — This annex supplements the preserved source artefact; it does not amend or replace original text or tables and is not approved policy or permission to operate.

Identity and evidence snapshot

- Use permanent AIR-ID for the registered system and stable AG-ID for each agent, paired together. Pre-registration discovery reference is provisional, not AIR-ID. Record exact ASBOM (AIG-AGT-04) version/revision and snapshot. Graph (AIG-AGT-05) is derived from ASBOM delegation edges and informs authority/reach evidence, not an independent source of authority.

Five dimensions versus AGPI and risk

- Record consequence, autonomy, authority, reach and controllability as a five-dimension profile with rationale/evidence for each, capability/multiplier flags, any mandatory tier floor, resulting agency tier and route. This is distinct from AGPI (AIG-ASS-01), which measures governance priority, and Risk Worksheet (AIG-ASS-02), which determines inherent/residual risk tiers. They answer different questions, are not substitutes, and none is approval.
- ASBOM (AIG-AGT-04) records capability, permissions, controls and delegation; its derived graph helps inspect authority paths/reach. A link, declared intention or graph alone does not prove a permission absent or enforced.

Re-entry and action evidence

- Agent-capable systems require full Intake (AIG-INV-03) and Agentic Triage (AIG-AGT-03); Fast Track (AIG-INV-02) cannot bypass review. Discovery of an unregistered system or material changes to purpose, model, data, access, tools, authority, agency, affected people or environment requires full triage re-entry and applicable AGPI/risk/specialist reassessment.
- AIG-AGT-06 Action ID/evidence links to AIR-ID, AG-ID and exact ASBOM version. Reference it in AIG-OPS-03/AIG-OPS-02/AIG-DEC-04 as applicable. Action evidence is not a gate decision, permission grant or risk acceptance.

Status of this proposed copy

Do not conflate capability map and agency profile

The proposed AIG-INV-05 Capability catalogue describes reusable service functions (UC-ID → CAP-ID → system AIR-ID); the five scored agentic dimensions and multipliers remain the separate AIG-AGT-02/AIG-AGT-03 agentic profile. Do not derive consequence, autonomy, authority, reach or controllability from a CAP-ID or a map edge. Assess and cite the underlying system/agent evidence, including the exact AIG-AGT-04 version; map coverage is not risk evidence or approval.

Proposed integrated draft; not approved and not effective. Original baseline content and tables are retained above. Formal approval, owner confirmation and document-control updates remain for the competent Council authority.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 51-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-040 — ISO/IEC 42001:2023 Artificial intelligence management system	Proposed management-system design mapping; licensed clause review and adoption needed.	https://www.iso.org/standard/42001
REQ-041 — ISO/IEC 23894:2023 AI — Guidance on risk management	Voluntary risk-management guidance; proposed alignment, not statutory conformity.	https://www.iso.org/standard/77304.html
NIST-AI-100-1 — NIST AI Risk Management Framework 1.0 (NIST AI 100-1, January 2023)	Voluntary alignment reference; not certification or statutory conformity	https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10