

PROPOSED INTEGRATED DRAFT — NOT APPROVED*For Council review only · Not authorised, effective or operational until formally approved*

AIG-AGT-02 · The Council · AI GOVERNANCE

Agentic Classification Reference

Apply the classification profile to the specific configuration and UC-ID/outcome being assessed, not only to the shared system. Distinct purposes, populations, action scopes or permissions may have distinct risk and delegated decisions under the same AIR-ID. UC-ID applies to ordinary uses too; action-capability classifications apply only when that configuration can take or initiate actions. A profile or map is not approval.

Capability vector, autonomy scale and agency tiers

Field	Controlled value
Document title	Agentic Classification Reference
Version	1.5 draft
Status	Draft for Council review; pre-adoption
Author	AI Governance Lead
Date	30 September 2026
Purpose	The shared vocabulary the Agent Record and the Agentic Triage both use, so a level or a capability means the same thing everywhere
Used by	Agent Record / ASBOM (Capability Vector and fields); Agentic Triage (agency profile and tiering)
Source	Playbook Appendix F (F.7.3, F.7.4); aligned to OWASP 2026 guidance (ASI03, ASI02, ASI08; Agent Control Standard, released 1 September 2026)
Classification	OFFICIAL
Change history	1.1 draft, 29 September 2026: QA remediation — applied findings QA-002, QA-004, QA-014, QA-033, QA-034, QA-053, QA-088, QA-120, QA-129, QA-143, QA-029 (see QA change log). Supersedes the unnumbered working draft of 19 September 2026. 1.2 draft, 29 September 2026: v3.7: the combination rule now uses the minimum risk-tier pathway each agency tier sets (AIG-AGT-03 §6; Gate Map AIG-DEC-01, Agentic pathway); permitted autonomy granted at Gate 6 (go-live); equivalent-control authorisation evidenced at Gate 2 and re-confirmed at Gate 6. Follow-up: “can it act?” gate wording replaced with “the ‘can it act?’ screen” / “screen-negative (cannot act)”; record value “Not applicable: cannot act”; T0 stated as action-capable (Gate 2, Gate 6 and the §4.4.6 floor apply). Proposed — for Council confirmation. 1.3 draft, 30 September 2026: v3.8: new section “Worked calibration examples” scores four realistic agents with Tables A and B (script output included) and states whether the results look proportionate; Tables A and B unchanged. The combination rule now lists the §4.4.4 impact floor (a confirmed Impact 5 sets at least Medium) with the §4.4.6 floors. Proposed — for Council confirmation. Final tidy-up: Example 1’s Low result now rests on the existing rule (a person reviews and sends each reply, so per-action review is evidenced and R3 does not apply), with the R3 exemption kept as an open question for the Council; Table C labels separate Consequence scores (Cons) from Table B rules (rule C1 to C5); script output regenerated with the same labels. 1.4 draft, 30 September 2026: v3.9.1: standards cross-reference note updated to the 54-entry requirement inventory (AIG-AIMS-05 REQ-

054, Civil Contingencies Act 2004, added).
 1.5 draft, 1 October 2026: v3.9.3: standards cross-reference note
 cites the 55-entry requirement inventory (AIG-AIMS-05 REQ-055,
 NIST AI 600-1 Generative AI Profile, added).

1. Why one reference

The Agent Record and the Agentic Triage both classify agents, and they must speak the same language. This reference fixes the definitions once. It has no logic of its own; it is the dictionary the other artefacts point to.

The organising idea, from which everything else follows: cognitive capability is not agency. An agent can be brilliant at reasoning and pose little danger, while a mediocre model wired to a production database, payments and email can pose a great deal. So we separate what an agent can think from what it can do, and we govern the doing.

2. Cognitive capability versus action capability

Keep the two apart when describing any agent.

- Cognitive (thinking): reason, plan, analyse, generate, retrieve, infer, recommend, decide. These raise quality and error questions.
- Action (doing): read, write, execute, communicate, purchase, modify, delegate, deploy. These raise agency and consequence questions, and are where governance effort belongs.

3. The capability vector

Record what the agent can technically do, capability by capability, as No, Yes or Scoped. Capability is not intention: an instruction not to do something is not a control if the ability remains.

Capability	Governance question
Read	What information can it access?
Write	What can it change?
Execute	What can it cause to happen?
Communicate	Who can it contact?
Purchase	Can it spend money?
Delegate	Can it create or subcontract another agent?
Persuade	Can it influence humans?
Code	Can it generate or execute software?
Discover	Can it independently find new resources?
Persist	Can it operate beyond a session?
Replicate	Can it create additional instances?
Learn	Can its behaviour change from experience?
Escalate	Can it request greater authority?

4. Agency multipliers

These capabilities change the risk picture sharply and are captured as explicit flags on the Agent Record, because any one of them can turn a modest agent into a high-agency one.

- The ten flags: persistence, memory, delegation, tool discovery, credential access, self-modification, replication, goal adaptation, external communication and financial authority.

5. The autonomy scale (A0 to A5)

Autonomy is permissioned, granted deliberately at Gate 6 (go-live), not inferred from how sophisticated the system looks. The level is recorded on the Agent Record and revisited at reauthorisation. Triage records only the assessed or requested level; the permitted level is granted by the go-live decision-maker at Gate 6 and recorded on the Agent Record with its decision reference (AIG-DEC-03).

Level	Meaning
A0 advisory	Agent recommends; a human acts.
A1 assisted	Agent prepares actions; a human explicitly executes.
A2 bounded execution	Agent executes pre-approved actions within strict limits.
A3 conditional autonomy	Agent selects actions within a defined policy envelope.
A4 delegated autonomy	Agent plans, selects tools, executes and adapts with limited human intervention.
A5 open-ended autonomy	Agent independently pursues complex objectives across changing environments.

Human oversight rule: per-action human authorisation is required before any irreversible or high-consequence action, at every autonomy level, unless an equivalent control is formally authorised at the relevant Council gate (evidenced at Gate 2; re-confirmed at Gate 6) and recorded in AIG-AGT-04. Post-action review and monitoring-only oversight are permitted only for non-consequential, reversible actions. Where it is unclear whether each action is reviewed by a person before it executes, “Unsure” is treated as “No”.

6. The agency tiers (T0 to T5)

The tier is the output of the Agentic Triage's agency profile, and it routes the agent to a governance pathway: each tier sets a minimum risk-tier pathway (AIG-AGT-03 §6; Gate Map AIG-DEC-01, Agentic pathway). It is based on actual behaviour and authority, not the marketing description. Tiers apply only to action-capable configurations, where the “can it act?” screen answer is Yes (“Unsure” is treated as Yes until confirmed); a screen-negative system (cannot act) receives no agency tier (record “Not applicable: cannot act”). T0 is still an action-capable tier: it sets no minimum pathway, but Gate 2 and Gate 6 apply and the §4.4.6 Critical floor applies (AIG-AGT-03 §6; AIG-DEC-01, Agentic pathway). The tier is assigned with the tier-assignment table below, not by reading these descriptions.

Tier	Meaning
T0 informational	Can act, but no action has a meaningful external effect: Consequence, Autonomy and Authority all 0, Reach and Controllability at most 1, and no multiplier flag.
T1 assisted	Low agency: Consequence, Autonomy and Authority at most 1 (each consequential action is individually executed or approved by a person), Reach and Controllability at most 2.
T2 bounded agent	Acts, but within tightly constrained authority.
T3 consequential agent	Can independently produce material external consequences.
T4 high-agency	High consequence (Consequence 4), difficult containment (Controllability 4), broad authority combined with significant autonomy or reach, or persistence, delegation or tool discovery combined with substantial autonomy or authority.

Tier	Meaning
T5 exceptional / high-consequence	Any dimension at 5; severe consequence with difficult containment (Consequence and Controllability both 4 or more); agent creation, replication or self-modification; or a mandatory escalation trigger.

7. The five agency dimensions

The triage scores these five and keeps them as a profile rather than adding them up, so a spike on one dimension is never hidden by low scores on the others. Controllability is reverse-scored: 5 means hardest to stop or reverse, because low controllability is the danger.

Dimension	Scale (0 to 5)
Consequence	0 trivial to 5 potentially catastrophic or irreversible. What happens if it gets it wrong?
Autonomy	0 acts only when a person triggers each individual action, to 5 open-ended autonomy. How independently can it act?
Authority	0 information only to 5 highly consequential authority. How much real-world power does it hold?
Reach	0 isolated to 5 broad, public or physical. How much of the environment can it affect?
Controllability	Reverse-scored: 0 instant stop with full rollback to 5 effectively irreversible (5 = hardest to stop or reverse). How easily can humans stop or reverse it?

Autonomy score and Autonomy Level: the 0–5 Autonomy score measures how independently the configuration can technically act; it is scored independently of the A0–A5 Autonomy Level, which is a permission. Persistence is recorded as a multiplier flag, not in the Autonomy score. Where the Autonomy score is higher than the number of the permitted A-level, capability exceeds intended authority: record the gap and show how the limit is technically enforced.

8. How they combine

Three rules govern how these pieces are used together.

- Profile, not sum: the five dimensions are read as a profile, not a sum. A system with modest autonomy but high authority and low controllability deserves attention a single averaged score would bury.
- Intensity: governance intensity scales with agency times capability times consequence times uncertainty. This is why the same agent can be low risk summarising documents and high risk moving money: the action and authority changed, not the agent.
- Combination with the risk tier: the governing pathway is the higher of (a) the Playbook §4.4 risk-tier pathway, including the §4.4.6 trigger floors and the §4.4.4 impact floor (a confirmed Impact score of 5 on any impact dimension sets at least Medium), and (b) the minimum risk-tier pathway set by the agency tier (T0 and T1 none, T2 Medium, T3 High, T4 High or Critical under §4.4.6, T5 Critical; AIG-AGT-03 §6 and the Gate Map AIG-DEC-01, Agentic pathway; Proposed — for Council confirmation); neither lowers the other.

9. Status and source

Draft for Council review, pre-adoption. The capability vector, the multipliers, the A0 to A5 scale, the T0 to T5 tiers and the five dimensions follow Playbook Appendix F.7.3 and F.7.4, adopted here as the common vocabulary and aligned to OWASP 2026 guidance (identity and privilege ASI03, tool misuse ASI02, cascading failures ASI08, and the Agent Control Standard, released 1 September 2026), verified against genai.owasp.org on 19 September 2026.

Draft for Council review. The dictionary the Agent Record and the Agentic Triage both point to.

Deterministic classification and escalation rules

Agency tier is a profile, not an average. The tier is set by the tier-assignment table below: the highest applicable floor governs (“highest floor wins”). The governance decision record must document the profile, the floors triggered and any escalation or exception.

- T0: the configuration can act, but no action has a meaningful external effect (every Table A floor is T0 and no multiplier flag is set).
- T1: low agency; the highest floor is T1, typically because each consequential action is individually executed or approved by a person.
- T2: bounded agent with explicit scoped execution and reliable containment.
- T3: consequential agent where actions can materially affect people, records, services, finances or external communications.
- T4: high-agency operation, as set by the T4 floors in Tables A and B: high consequence, difficult containment, broad authority with significant autonomy or reach, or persistence, delegation or tool discovery with substantial autonomy or authority.
- T5: exceptional/high-consequence agency, as set by the T5 floors: any dimension at 5, severe consequence with difficult containment, agent creation, replication or self-modification, or a mandatory escalation trigger.

Tier-assignment table — Proposed — for Council confirmation

PROPOSED — FOR COUNCIL CONFIRMATION. This table makes the agentic triage deterministic. Score each of the five dimensions 0–5 with rationale and evidence, record the ten multiplier flags, apply every rule in Tables A and B, and assign the highest floor triggered (minimum T0). In Playbook Appendix F.7.4 terms, each cell of Table A and each rule in Table B is a defined escalation threshold. An unknown or unevidenced score is scored at the higher plausible value, and an unknown flag is treated as present (fail safe).

Table A — per-dimension floors (score → minimum agency tier; Controllability 5 = hardest to stop or reverse)

Score	Consequence	Autonomy	Authority	Reach	Controllability
0	T0	T0	T0	T0	T0
1	T1	T1	T1	T0	T0
2	T2	T2	T2	T1	T1
3	T3	T2	T3	T2	T2
4	T4	T3	T3	T3	T4
5	T5	T5	T5	T5	T5

Table B — combination, multiplier-flag and trigger floors

Rule	Condition	Tier floor
B1	Two or more of Autonomy, Authority and Reach score 4 or more	T4
B2	Consequence 4 or more and Controllability 4 or more	T5
C1	Memory or external communication flag	T1
C2	Any one of the persistence, delegation, tool discovery, credential access or goal adaptation flags	T2
C3	Financial authority flag	T3
C4	Persistence, delegation or tool discovery flag, together with Autonomy 3 or more or Authority 3 or more	T4
C5	Replication or self-modification flag, or the agent can create new agents	T5
D1	Any mandatory escalation trigger (listed below) is present	T5: executive and safety escalation; no

Rule	Condition	Tier floor
		consequential production use until resolved

Tie-break and recording: the highest floor wins. The computed tier is the highest floor triggered. No rule, average, score or adjustment lowers a floor: an adjustment may only raise the computed tier. Record each score with its rationale and evidence, the flags, every floor triggered, the resulting tier and the rule(s) that set it, on the Agent Record (AIG-AGT-04) and in the governance case. Any upward adjustment is recorded with its reason and authority basis (the officer or forum with confirmed delegation) in AIG-DEC-03, and the computed tier is always kept alongside it.

Assurance check: the table was tested by enumerating all 7,776 possible five-dimension profiles (each dimension scored 0–5), and again under every combination of the ten multiplier flags. Every profile maps to exactly one tier, raising any score never lowers the tier, and the T0 and T1 criteria match the table.

Combination with the risk tier: the governing pathway is the higher of (a) the Playbook §4.4 risk-tier pathway, including the §4.4.6 trigger floors and the §4.4.4 impact floor (a confirmed Impact score of 5 on any impact dimension sets at least Medium), and (b) the minimum risk-tier pathway set by the agency tier (T0 and T1 none, T2 Medium, T3 High, T4 High or Critical under §4.4.6, T5 Critical; AIG-AGT-03 §6 and the Gate Map AIG-DEC-01, Agentic pathway; Proposed — for Council confirmation); neither lowers the other.

Mandatory escalation triggers (Table B, rule D1) include: authority ceiling cannot be reliably enforced; containment cannot be demonstrated; delegation/recursion is unbounded; irreversible high-consequence actions are possible without appropriate human control; persistent/shared memory lacks effective isolation; or the agent can create or modify agents/configurations without separately authorised controls.

Frontier model is recorded separately from agency tier and never substitutes for consequence assessment.

Worked calibration examples (Proposed — for Council confirmation)

PROPOSED — FOR COUNCIL CONFIRMATION. Four realistic agents scored with the tier-assignment table exactly as printed above (Tables A and B and the multiplier flags), to show what the rules produce in practice. Scores follow the scale anchors in section 7; each has a one-line rationale. The agency tier gives the agency minimum pathway (AIG-AGT-03 §6), which is combined with an assumed Playbook §4.4 risk tier, the §4.4.4 impact floor and any §4.4.6 trigger floors: the governing tier is the highest of these. Risk scores are illustrative assumptions, not assessments of any real system. The results were computed with a short script that implements Tables A and B cell for cell (checked against all 7,776 profiles for the T0 and T1 criteria); its output is reproduced below.

Example 1 — Email-drafting copilot. Drafts replies in an officer's mailbox; the officer edits and sends each one.

- Consequence 1: a poor draft could mislead a resident only if the officer sends it unchanged; a person reads and sends every reply.
- Autonomy 1: a draft is prepared for each incoming message without a separate request, but it cannot send, file or delete.
- Authority 1: it can create draft items only; no sending and no changes to other records.
- Reach 1: one officer's mailbox.
- Controllability 0: a draft can be deleted at once and nothing leaves the mailbox until a person sends it.
- Flags: none (no memory across sessions; no external communication, as it cannot send; it works in the signed-in officer's session and holds no credentials of its own).
- Assumed risk: internal productivity, Likelihood 2 × Impact 2 = 4 (Low). No §4.4.6 trigger: the copilot's own action is creating a draft for a person, and the officer reviews and sends each reply, so per-action

human review is evidenced and rule R3 does not apply. This rests on the existing rule, not on any exemption for low-consequence or reversible actions.

Example 2 — Resident booking agent. A public chat agent that books and cancels appointments (for example bulky-waste collections) in one booking system through its API.

- Consequence 2: a wrong booking or cancellation causes a missed or delayed appointment; the harm is recoverable.
- Autonomy 2: executes book and cancel actions within fixed limits once the resident confirms; no staff involvement.
- Authority 2: create and cancel bookings in one system; no payments and no case records.
- Reach 3: public-facing and used by many residents, through one back-end system.
- Controllability 2: can be switched off at the API; a cancelled slot can usually be reinstated but may already have been re-let.
- Flags: credential access (holds an API service credential) and external communication (sends confirmations to residents).
- Assumed risk: Likelihood 3 × Impact 3 = 9 (Medium). §4.4.6 trigger: bookings and cancellations execute without Council review of each action, and the resident's confirmation is not recorded as per-action human authorisation, so rule R3 sets the Critical floor ("Unsure" is treated as "No").

Example 3 — Housing case-triage agent. Reads new housing and homelessness case files and writes a priority flag to each case in the case-management system; an officer reviews the case before any decision.

- Consequence 4: a wrongly low flag can delay help to a household at risk of homelessness; the harm can be serious and hard to undo.
- Autonomy 3: runs on every new case and chooses the flag within a defined policy, without a request per case.
- Authority 2: writes one field (the priority flag) in one system; it cannot decide, close or message.
- Reach 3: the whole housing caseload, including vulnerable residents, in one internal system.
- Controllability 2: a flag can be changed or the agent stopped at once, but a wrong flag may go unnoticed until the case is reached.
- Flags: credential access (service account) and persistence (runs continuously, beyond a session).
- Assumed risk: Likelihood 3 × Impact 4 = 12 (High). §4.4.6 triggers: housing and homelessness decisions and vulnerable residents (High floor); flags are written without review of each flag, so rule R3 sets the Critical floor.

Example 4 — Multi-agent decision-letter workflow. An orchestrating agent directs a fixed set of sub-agents that gather evidence from several systems, draft decision letters and send them; staff review a sample after sending.

- Consequence 4: an incorrect decision letter can wrongly refuse or reduce an entitlement; correction needs a review or appeal.
- Autonomy 4: plans the evidence gathering, selects tools and sends with limited human intervention.
- Authority 4: issues formal decisions to residents on the Council's behalf.
- Reach 4: many residents and several internal systems.
- Controllability 4: a sent letter cannot be recalled; correcting it needs a new decision.
- Flags: delegation (orchestrator to sub-agents), memory (shared case memory), persistence, credential access and external communication. Mandatory escalation trigger (rule D1): irreversible high-consequence actions are possible without appropriate human control, because only a sample is reviewed.
- Assumed risk: Likelihood 3 × Impact 5 = 15 (High); the §4.4.4 impact floor (Medium) is met but lower; §4.4.6 triggers: statutory decisions and autonomous action without per-action review (Critical).

Table C — calibration results (illustrative; scores: Cons = Consequence, Au = Autonomy, Ah = Authority, Reach, Ct = Controllability; Table B rules are written “rule” plus their ID, for example rule C4; rule R3 is the §4.4.6 trigger for actions without per-action review)

Example	Scores (Cons / Au / Ah / Reach / Ct) and flags	Score floors (Table A) and rules (Table B) → agency tier	Agency minimum pathway	Assumed risk tier and trigger floors	Governing tier
1 Email-drafting copilot	1 / 1 / 1 / 1 / 0; flags: none	Cons 1, Au 1, Ah 1 → T1 (set by Cons 1, Au 1, Ah 1)	None	Low (4); trigger floors: none	Low
2 Resident booking agent	2 / 2 / 2 / 3 / 2; flags: credential access, external communication	Cons 2, Au 2, Ah 2, Reach 3, Ct 2; rules C1, C2 → T2 (set by Cons 2, Au 2, Ah 2, Reach 3, rule C2)	Medium	Medium (9); trigger floors: Critical (rule R3, no per-action review)	Critical
3 Housing case-triage agent	4 / 3 / 2 / 3 / 2; flags: credential access, persistence	Cons 4, Au 3, Ah 2, Reach 3, Ct 2; rules C2, C4 → T4 (set by Cons 4, rule C4)	Critical	High (12); trigger floors: High (housing, vulnerable residents); Critical (rule R3, no per-action review)	Critical
4 Multi-agent decision-letter workflow	4 / 4 / 4 / 4 / 4; flags: credential access, delegation, external communication, memory, persistence	Cons 4, Au 4, Ah 4, Reach 4, Ct 4; rules B1, B2, C1, C2, C4, D1 → T5 (set by rules B2, D1)	Critical	High (15); impact floor Medium; trigger floors: Critical (statutory decisions); Critical (rule R3, no per-action review)	Critical

Script output. Produced by the calibration script from the scores and flags above:

Ex1 Email-drafting copilot: scores Cons 1, Au 1, Ah 1, Reach 1, Ct 0; flags none; floors A:Consequence 1, A:Autonomy 1, A:Authority 1; agency tier T1 (set by A:Consequence 1, A:Autonomy 1, A:Authority 1); agency minimum none; risk $L2 \times I2=4$ Low; no trigger floor → governing tier Low

Ex2 Resident booking agent: scores Cons 2, Au 2, Ah 2, Reach 3, Ct 2; flags credential access, external communication; floors A:Consequence 2, A:Autonomy 2, A:Authority 2, A:Reach 3, A:Controllability 2, rule C1, rule C2; agency tier T2 (set by A:Consequence 2, A:Autonomy 2, A:Authority 2, A:Reach 3, rule C2); agency minimum Medium; risk $L3 \times I3=9$ Medium; trigger floors Critical → governing tier Critical

Ex3 Housing case-triage agent: scores Cons 4, Au 3, Ah 2, Reach 3, Ct 2; flags credential access, persistence; floors A:Consequence 4, A:Autonomy 3, A:Authority 2, A:Reach 3, A:Controllability 2, rule C2, rule C4; agency tier T4 (set by A:Consequence 4, rule C4); agency minimum Critical; risk $L3 \times I4=12$ High; trigger floors High, Critical → governing tier Critical

Ex4 Multi-agent decision-letter workflow: scores Cons 4, Au 4, Ah 4, Reach 4, Ct 4; flags credential access, delegation, external communication, memory, persistence; floors A:Consequence 4, A:Autonomy 4, A:Authority 4, A:Reach 4, A:Controllability 4, rule B1, rule B2, rule C1, rule C2, rule C4, rule D1; agency tier T5 (set by rule B2, rule D1); agency minimum Critical; risk $L3 \times I5=15$ High; impact floor Medium; trigger floors Critical, Critical → governing tier Critical

Sensitivity Ex2: with per-action review evidenced (or an equivalent control authorised), R3 does not apply → governing tier Medium

Sensitivity Ex3: if an officer confirmed each flag before it is written, R3 does not apply → governing tier High

What this tells the Council. Two of the four results look proportionate and two raise a calibration question. The drafting copilot (T1, governing tier Low) and the decision-letter workflow (T5, Critical with executive and safety escalation) land where most practitioners would expect. The copilot's Low result follows from the rules as written, not from an exemption: its only action is creating a draft for a person, and a person reviews and sends each reply, so per-action human review is evidenced and R3 does not apply. The housing triage agent's Critical result is defensible, given vulnerable households and flags written without review of each one, although it rests on rule R3 rather than on the tier table: if an officer confirmed each flag, the result would be High. The booking agent is the outlier. The tier table places it at T2 (minimum Medium) and its assumed risk tier is Medium, but rule R3 (the §4.4.6 trigger for actions executed without per-action human review) lifts it to Critical. R3 is the bluntest rule in the set: it takes no account of the Consequence or Controllability scores, so a low-consequence, largely reversible booking agent receives the same pathway as an agent issuing statutory decisions, and the T2 and T3 rows of the pathway table are in effect bypassed for any agent that acts without per-action review. The next most likely to over-classify is "any single dimension at 5 → T5" (Table A, score 5): it did not bite here, but a public information agent scored Reach 5 only because it is public-facing would go straight to T5 and Critical. The credential-access flag (C2) places every API-integrated agent at T2 or above, which is cautious but consistent with the fail-safe design. This section does not change Tables A and B. Questions for the Council: whether R3 should exempt low-consequence, reversible actions such as the bookings and cancellations in Example 2 (for example by applying only where Consequence is 3 or more or the action is irreversible, in line with the post-action review allowance in the human oversight rule above), and whether a resident's explicit confirmation of each action can count as per-action human authorisation when recorded in AIG-AGT-04.

Annex — Classification identity and routing boundaries

PROPOSED INTEGRATION ONLY — This annex supplements the preserved source artefact; it does not amend or replace original text or tables and is not approved policy or permission to operate.

Identity and evidence snapshot

- Use permanent AIR-ID for the registered system and stable AG-ID for each agent, paired together. Pre-registration discovery reference is provisional, not AIR-ID. Record exact ASBOM (AIG-AGT-04) version/revision and snapshot. Graph (AIG-AGT-05) is derived from ASBOM delegation edges and informs authority/reach evidence, not an independent source of authority.

Five dimensions versus AGPI and risk

- Record consequence, autonomy, authority, reach and controllability as a five-dimension profile with rationale/evidence for each, capability/multiplier flags, any mandatory tier floor, resulting agency tier and route. This is distinct from AGPI (AIG-ASS-01), which measures governance priority, and Risk Worksheet (AIG-ASS-02), which determines inherent/residual risk tiers. They answer different questions, are not substitutes, and none is approval.
- ASBOM (AIG-AGT-04) records capability, permissions, controls and delegation; its derived graph helps inspect authority paths/reach. A link, declared intention or graph alone does not prove a permission absent or enforced.

Re-entry and action evidence

- Agent-capable systems require full Intake (AIG-INV-03) and Agentic Triage (AIG-AGT-03); Fast Track (AIG-INV-02) cannot bypass review. Discovery of an unregistered system or material changes to purpose, model, data, access, tools, authority, agency, affected people or environment requires full triage re-entry and applicable AGPI/risk/specialist reassessment.
- AIG-AGT-06 Action ID/evidence links to AIR-ID, AG-ID and exact ASBOM version. Reference it in AIG-OPS-03/AIG-OPS-02/AIG-DEC-04 as applicable. Action evidence is not a gate decision, permission grant or risk acceptance.

Status of this proposed copy

Do not conflate capability map and agency profile

The proposed AIG-INV-05 Capability catalogue describes reusable service functions (UC-ID → CAP-ID → system AIR-ID); the five scored agentic dimensions and multipliers remain the separate AIG-AGT-02/AIG-AGT-03 agentic profile. Do not derive consequence, autonomy, authority, reach or controllability from a CAP-ID or a map edge. Assess and cite the underlying system/agent evidence, including the exact AIG-AGT-04 version; map coverage is not risk evidence or approval.

Proposed integrated draft; not approved and not effective. Original baseline content and tables are retained above. Formal approval, owner confirmation and document-control updates remain for the competent Council authority.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 55-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-040 — ISO/IEC 42001:2023 Artificial intelligence management system	Proposed management-system design mapping; licensed clause review and adoption needed.	https://www.iso.org/standard/42001
REQ-041 — ISO/IEC 23894:2023 AI — Guidance on risk management	Voluntary risk-management guidance; proposed alignment, not statutory conformity.	https://www.iso.org/standard/77304.html
NIST-AI-100-1 — NIST AI Risk Management Framework 1.0 (NIST AI 100-1, January 2023)	Voluntary alignment reference; not certification or statutory conformity	https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10