

PROPOSED INTEGRATED DRAFT — NOT APPROVED*For Council review only · Not authorised, effective or operational until formally approved*

AIG-AGT-03 · The Council · AI GOVERNANCE

Agentic Triage

Function specification for the two-function triage route


Field	Controlled value
Document title	Agentic Triage: Function Specification
Version	1.6 draft
Status	Draft for Council review; pre-adoption
Author	AI Governance Lead
Date	30 September 2026
Purpose	Specify how the triage route handles agents: one front door, two functions, dynamic assessment
Built to	Agentic AI Governance Design Principles (AIG-AGT-01); implements Playbook Appendix F.2 and F.7.4
Uses / writes	Agentic Classification Reference (vocabulary and tier-assignment table); writes the assessed Agency Tier and the assessed or requested Autonomy Level to the Agent Record / ASBOM (the permitted Autonomy Level is recorded only after the Gate 6 (go-live) decision)
Source	OWASP Top 10 for Agentic Applications 2026 and Agent Control Standard (released 1 September 2026); Playbook Appendix F.2 and F.7
Classification	OFFICIAL
Change history	1.1 draft, 29 September 2026: QA remediation — applied findings QA-002, QA-004, QA-014, QA-016, QA-053, QA-057, QA-088, QA-120, QA-157, QA-171, QA-201, QA-029 (see QA change log). Supersedes the unnumbered working draft of 19 September 2026. 1.2 draft, 29 September 2026: v3.7: §6 pathway table gains the minimum risk-tier pathway per agency tier (the Gate Map, AIG-DEC-01, now carries these pathways in its Agentic pathway section); Gate 2 mandatory for action-capable uses; permitted autonomy granted at Gate 6; “control gates” renamed agentic control checkpoints and mapped to Gate 2 and Gate 6. Follow-up: R1–R3 apply to every action-capable use (any agency tier, including T0; Unsure treated as Yes); T0 row: no minimum pathway and no further agentic triage, Gate 2 (agentic control checkpoints), Gate 6 (permitted autonomy) and the §4.4.6 floor still apply; “Reading the figure” text aligned; “the gate” wording for the can-it-act screen replaced with “the ‘can it act?’ screen”. v3.8 (30 September 2026, no version change): the combination rule (opening paragraph, “Reading the figure” and §6) lists the §4.4.4 impact floor, and the opening paragraph says the AGPI priority sets urgency only, not the route. Proposed — for Council confirmation. 1.3 draft, 30 September 2026: v3.9.1: standards cross-reference note updated to the 54-entry requirement inventory (AIG-AIMS-05 REQ-054, Civil Contingencies Act 2004, added). 1.4 draft, 1 October 2026: v3.9.3: standards cross-reference note cites the 55-entry requirement inventory (AIG-AIMS-05 REQ-055, NIST AI 600-1 Generative AI Profile, added). 1.5 draft, 2 October 2026: v3.9.5: standards cross-reference note

cites the 56-entry requirement inventory (AIG-AIMS-05 REQ-056, SI 2026/425 ICO statutory code of practice on AI and automated decision-making, added). Proposed — for Council confirmation.
1.6 draft, 4 October 2026: v3.9.9: §6 pathway rows and the rules for action-capable uses name the Agent Record level (core; plus feature sheets; full ASBOM at T3 and above; Playbook F.3). Proposed — for Council confirmation.

1. What this specifies

The agentic layer keeps one admission route and adds a distinct, dynamic agentic triage for agents (Playbook Appendix F.2). This document specifies how that works: the “can it act?” screen, the two functions, the steps inside the agentic function, the output and how it connects to the rest of the layer.

2. The route at a glance



Action risk is re-evaluated at runtime; a system moves between risk states as its action and authority change.

One front door and one AIR-ID. Every AI system is priority-triaged; anything that can act is additionally agency-triaged; the AGPI priority sets urgency and sequencing only, and the route follows the governing pathway: the higher of the Playbook §4.4 risk-tier pathway (including the §4.4.6 trigger floors and the §4.4.4 impact floor) and the agency-tier minimum (see §6).

Reading the figure: a “No” answer at the “can it act?” screen goes to the standard AI review pathway (the normal route), not to function two. A “Yes” (or “Unsure”, treated as Yes until confirmed) makes the use action-capable and runs the agentic triage: five dimensions give an agency tier (T0 to T5), and the tier sets a minimum risk-tier pathway (T0 and T1 none; T2 Medium; T3 High; T4 High, or Critical where actions run without evidenced per-action human review; T5 Critical plus executive and safety escalation). The governing pathway is the higher of the risk-tier pathway (Playbook §4.4, including the §4.4.6 trigger floors and the §4.4.4 impact floor (a confirmed Impact score of 5 sets at least Medium)) and this minimum; no new board is added. The existing Council gates then apply: for every action-

capable use (any agency tier, including T0), Gate 2 (technical design) is mandatory, Gate 6 (go-live) grants the permitted autonomy level and the §4.4.6 Critical floor applies where actions run without evidenced per-action human review; the permitted autonomy level is recorded in the decision record (AIG-DEC-03) and written to the Agent Record (AIG-AGT-04). The full table is in §6 and in the Gate Map (AIG-DEC-01, Agentic pathway).

Run the system intake once under AIR-ID, then identify each materially distinct use/action with its own UC-ID. Determine priority, risk, route and delegated decision scope per use; do not let a lower-risk system baseline suppress a higher-risk UC. Only an actually action-capable configuration receives the additional agentic/action/authority triage; ordinary tools remain non-agentic while still receiving UC-ID scope.

3. Function one: system and priority triage

This is the existing triage, run for every AI system, unchanged except for one addition.

- Priority and risk: AGPI scores governance priority and the risk worksheet sets inherent and residual risk tiers, exactly as now.
- The “can it act?” screen (not one of the Council decision gates, Gate 0 to Gate 8): one question is added to function one: can it act? It is asked first as the screen question at intake (AIG-INV-02 Part B; AIG-INV-03), so Light-touch cases are caught too, and it is confirmed after priority and risk assessment (Playbook Appendix F.2). If no, the system continues down the standard AI review path and the agentic function does not fire. If yes (or unsure, which is treated as yes until confirmed), the use is action-capable and function two runs. For each action-capable UC-ID also record: is each action reviewed by a person before it executes? Answer “Yes” only where evidenced; “Unsure” is treated as “No”, which engages the Playbook §4.4.6 Critical floor for that UC-ID.

4. Function two: the agentic triage

Fires only for agents. It runs in this order, and it governs the authority envelope, what the agent can do, not what it is meant to do.

- 1 Freeze the system as an object: complete the Agentic System Snapshot (purpose, owner, environment, models, agents, memory, tools, data, interfaces, human involvement, persistence, delegation, authority). Establish what exists before scoring anything.
- 2 Map the action chain and mark every point where the system can change reality, so governance effort lands where consequences occur.
- 3 Separate cognitive capability from action capability, and record the capability vector. A clever model with little agency is not the concern; a modest model wired to payments and email is.
- 4 Establish the authority boundary: the maximum the agent can do without asking anyone. Look for capability exceeding intended authority.
- 5 Score the five agency dimensions in the Agentic Classification Reference (AIG-AGT-02): consequence, autonomy, authority, reach and controllability. Record a 0–5 profile, not a sum or average. Controllability is reverse-scored: 5 means hardest to stop or reverse. Flag persistence, delegation and the other agency multipliers separately.
- 6 Run the worst plausible chain: starting from a legitimate task, the most consequential chain of actions the architecture permits.
- 7 Test the boundaries (authority, tool, data, identity, delegation, memory, instruction, escalation, persistence, termination), not just the model.
- 8 Run the kill-switch test: not could you stop it, but show me. Demonstrate stopping it, revoking credentials, disabling tools, handling actions in flight, and reconstructing what it did.
- 9 Assign the agency tier T0 to T5 with the tier-assignment table (§6, reproduced from AIG-AGT-02): the highest floor wins. Record the scores, flags, floors triggered, resulting tier and the rule(s) that set it. Record the assessed or requested autonomy level only: the permitted autonomy level is granted at

Gate 6 (go-live) by the go-live decision-maker, recorded in AIG-DEC-03 and written to the Agent Record (AIG-AGT-04).

5. Dynamic, not one-time: the three linked assessments

Assessment join — retain the system AIR-ID and cite the applicable UC-ID in each risk, authority, gate, action and incident record. Shared system evidence may be reused only with an explicit applicability check; use-specific outcome, action scope, residual risk and decision must remain separately traceable.

A single classification at deployment is treated as insufficient, because the same agent moves between risk states. So the agentic function produces three linked assessments rather than one score.

- System risk: what could the overall system cause? Set for the gate decisions (Gate 2 and Gate 6) and at reauthorisation.
- Action risk: how risky is this particular action? Evaluated per consequential action at runtime.
- Authority risk: how much power has been delegated to it? Read from the Authority Graph and revisited when authority changes.

The worked case: an HR agent is low risk summarising CVs, higher ranking candidates, and high autonomously rejecting them. The agent has not changed; its action and authority have. Action risk is where that is caught, which is why it is a runtime function, not a gate score.

6. Output and routing

The tier routes the agent to a governance pathway. Because the five dimensions form a profile, a high-consequence or hard-to-contain reading can escalate the pathway regardless of the other scores. On the reverse-scored controllability scale, harder to contain means a higher score (5 = hardest to stop or reverse). The tier is assigned with the table below, reproduced from the Agentic Classification Reference (AIG-AGT-02); if the two ever differ, AIG-AGT-02 governs.

Tier-assignment table — Proposed — for Council confirmation

PROPOSED — FOR COUNCIL CONFIRMATION. This table makes the agentic triage deterministic. Score each of the five dimensions 0–5 with rationale and evidence, record the ten multiplier flags, apply every rule in Tables A and B, and assign the highest floor triggered (minimum T0). In Playbook Appendix F.7.4 terms, each cell of Table A and each rule in Table B is a defined escalation threshold. An unknown or unevidenced score is scored at the higher plausible value, and an unknown flag is treated as present (fail safe).

Table A — per-dimension floors (score → minimum agency tier; Controllability 5 = hardest to stop or reverse)

Score	Consequence	Autonomy	Authority	Reach	Controllability
0	T0	T0	T0	T0	T0
1	T1	T1	T1	T0	T0
2	T2	T2	T2	T1	T1
3	T3	T2	T3	T2	T2
4	T4	T3	T3	T3	T4
5	T5	T5	T5	T5	T5

Table B — combination, multiplier-flag and trigger floors

Rule	Condition	Tier floor
B1	Two or more of Autonomy, Authority and Reach score 4 or more	T4
B2	Consequence 4 or more and Controllability 4 or more	T5
C1	Memory or external communication flag	T1
C2	Any one of the persistence, delegation, tool discovery, credential access or goal adaptation flags	T2
C3	Financial authority flag	T3
C4	Persistence, delegation or tool discovery flag, together with Autonomy 3 or more or Authority 3 or more	T4
C5	Replication or self-modification flag, or the agent can create new agents	T5
D1	Any mandatory escalation trigger (listed below) is present	T5: executive and safety escalation; no consequential production use until resolved

Tie-break and recording: the highest floor wins. The computed tier is the highest floor triggered. No rule, average, score or adjustment lowers a floor: an adjustment may only raise the computed tier. Record each score with its rationale and evidence, the flags, every floor triggered, the resulting tier and the rule(s) that set it, on the Agent Record (AIG-AGT-04) and in the governance case. Any upward adjustment is recorded with its reason and authority basis (the officer or forum with confirmed delegation) in AIG-DEC-03, and the computed tier is always kept alongside it.

Assurance check: the table was tested by enumerating all 7,776 possible five-dimension profiles (each dimension scored 0–5), and again under every combination of the ten multiplier flags. Every profile maps to exactly one tier, raising any score never lowers the tier, and the T0 and T1 criteria match the table.

Combining the two readings: each agency tier sets a minimum risk-tier pathway (table below; Proposed — for Council confirmation). The governing pathway is the higher of (a) the Playbook §4.4 risk-tier pathway, including the §4.4.6 trigger floors and the §4.4.4 impact floor (a confirmed Impact score of 5 sets at least Medium), and (b) that minimum; neither lowers the other. No new board is created: the existing forums, gates and records in the Gate Map (AIG-DEC-01) apply for the governing pathway.

Agency tier	Minimum risk-tier pathway	Agentic requirement	Governing pathway (higher governs)
T0 informational	None (no minimum pathway)	No minimum pathway and no further agentic triage; as an action-capable use, Gate 2 (agentic control checkpoints) and Gate 6 (permitted autonomy) still apply (R1, R2), and the §4.4.6 floor applies (R3); recorded on the Agent Record (AIG-AGT-04; core record)	The Playbook §4.4 risk-tier pathway; Gate 2 and Gate 6 apply; the §4.4.6 Critical floor applies to any UC-ID whose actions execute without per-action human review
T1 assisted	None	Normal route (standard AI review); Gate 2 and Gate 6 mandatory; recorded on the Agent Record (AIG-AGT-04; core record, plus the sheets its features need)	Higher of the minimum and the Playbook §4.4 risk-tier pathway; Gate 2 and Gate 6 apply; the §4.4.6 Critical floor applies to any UC-ID whose actions execute without per-

Agency tier	Minimum risk-tier pathway	Agentic requirement	Governing pathway (higher governs)
			action human review
T2 bounded agent	Medium	Agentic controls review at Gate 2 (Technical design authority, with AI Governance Working Group input); the forum for the governing pathway decides	Higher of Medium and §4.4; the §4.4.6 Critical floor applies to any UC-ID whose actions execute without per-action human review
T3 consequential agent	High	As T2, plus a formal AI Assurance Board recommendation before the decision	Higher of High and §4.4; the §4.4.6 Critical floor applies to any UC-ID whose actions execute without per-action human review
T4 high-agency	High (Critical where actions run without evidenced per-action human review — §4.4.6)	High-agency review: as T3, and the AIG-ASS-11 security review (ASI01–ASI10) complete before the Gate 2 decision	Higher of High and §4.4; Critical where the §4.4.6 floor applies
T5 exceptional / high-consequence	Critical	Executive and safety escalation (SRO and Executive Sponsor oversight; independent challenge); no consequential production use until resolved	Critical, with executive and safety escalation in every case

Rules for every action-capable use (any agency tier, including T0); “action-capable” means the “can it act?” answer is Yes, with “Unsure” treated as Yes until confirmed. R1: Gate 2 (Technical design review) is mandatory; the agentic control checkpoints (listed under Minimum agentic control checkpoints below) are evidenced there and re-confirmed at Gate 6. R2: Gate 6 (go-live) grants the permitted autonomy level for each UC-ID; the go-live decision-maker records it in AIG-DEC-03 and it is written to AIG-AGT-04; it may be lower than requested. R3: the §4.4.6 Critical floor (actions without evidenced per-action human review; “Unsure” is treated as “No”) is unchanged and applies to every action-capable use, T0 included. Record level: every action-capable use has at least the core Agent Record; persistence, memory, tool discovery or credential access and delegation add their own sheets; T3 and above, or financial authority, needs the full ASBOM (Playbook F.3; AIG-AGT-04 column BL). R4: agentic triage is a routing and evidence mechanism, not an approval layer; decisions stay with the officer or forum holding confirmed delegation (AIG-DEC-03 or approved minutes), with dated events in AIG-DEC-04.

7. How it connects to the rest of the layer

- One record: everything enters the same route and keeps one AIR-ID; the priority reading comes from AGPI on the Register (AIG-INV-04).
- Writes to the ASBOM: the agentic function writes the assessed Agency Tier and the assessed or requested Autonomy Level onto the Agent Record / ASBOM (the permitted Autonomy Level is added only after the Gate 6 (go-live) decision, with its AIG-DEC-03 reference), and its boundary and kill-switch results become evidence there.
- Routing: the Gate Map (AIG-DEC-01) now carries these pathways (its “Agentic pathway” section); the routing is extended, not replaced.

8. Deployment-time now, runtime next

Consistent with the design principles (AIG-AGT-01) and Playbook Appendix F.7.8, the function is built in two parts.

- Buildable now: the “can it act?” screen, the snapshot, the profile and multipliers, the worst-plausible-chain, the boundary and kill-switch checks, the tiering and the routing.
- Production requirement (checked at Gate 6): per-action risk evaluation and continuous authorisation are required for consequential autonomy where the approved authority envelope depends on them. Later platform delivery is a design roadmap, not evidence or permission to operate. A design-only or supervised pilot may test proposed controls only within the exact UC-ID scope authorised by a delegated decision; it is not production approval. A system-level move to Approved/Active is not a production decision. Consequential autonomy may operate for a UC-ID only when each required runtime control is Implemented and Evidenced, or a time-bounded compensating control has been expressly accepted for that use by the relevant delegated authority, with owner, expiry, residual risk and evidence recorded in AIG-DEC-03 (decision) and AIG-DEC-04 (gate event and condition). If neither exists, keep that UC-ID outside consequential production use.

9. The first-meeting checklist

For the first governance conversation about any agent, these ten questions are answered before moving on. They are the human-readable form of the agentic triage.

- 1 What is the agent's purpose?
- 2 What can it actually do, not what was it intended to do?
- 3 What systems and data can it access?
- 4 What can it change in the real world?
- 5 What authority has been delegated to it?
- 6 Can it act without a human?
- 7 Can it persist, remember, delegate or create additional actors?
- 8 What is the worst plausible chain of actions?
- 9 Can we stop it immediately and prove what it did?
- 10 Who is accountable when it causes a consequential outcome?

10. Status and source

Draft for Council review, pre-adoption; nothing evidenced until operated. The method (snapshot, action chain, five dimensions, multipliers, boundary and kill-switch tests, T0 to T5, three linked assessments, the ten questions) follows Playbook Appendix F.2 and F.7, adopted to satisfy OWASP 2026 controls: identity and privilege (ASI03), tool misuse (ASI02), cascading failures (ASI08), human-agent trust (ASI09) and rogue agents (ASI10), with the dynamic stance from the Agent Control Standard (released 1 September 2026), verified against genai.owasp.org on 19 September 2026.

Draft for Council review. Two functions, one route, one register; writes to the Agent Record / ASBOM. Operational decision rules

The Agentic Triage Function must produce a reproducible AIR-ID-linked assessment. The assessor records a 0–5 score, rationale and evidence for each of the same five dimensions: consequence, autonomy, authority, reach and controllability. Persistence, delegation and other agency multipliers are recorded separately as capability flags and control triggers; they do not replace autonomy or authority scores. Apply the tier-assignment table (§6; AIG-AGT-02), where the highest applicable floor wins, record the resulting route and required controls, and write the assessed Agency Tier and the assessed or requested Autonomy Level to the Agent Record / ASBOM (AIG-AGT-04). The permitted Autonomy Level is granted at Gate 6 (go-live) by the go-live decision-maker and recorded only after that decision, with its AIG-DEC-03 reference. Retain the scored profile and evidence reference with the governance case.

Use AIG-ASS-11 to record proportionate ASI01–ASI10 applicability and tests for action-capable agents at any base risk tier, and AIG-AGT-04 for each required runtime control state and evidence. Set AIG-OPS-02 monitoring thresholds before go-live (thresholds agreed for the decision; live only after Gate 6); AIG-AGT-06 records consequential actions. AI Assurance Board views are assurance opinions; the decision is made under the Council delegation.

Minimum agentic control checkpoints

These are agentic control checkpoints, not Council decision gates. Each checkpoint's evidence is received at Gate 2 (Technical design review) and re-confirmed at Gate 6 (go-live); the checkpoints are memory control, tool/interface authority, delegation chain, containment/kill switch, identity/credentials and rollback.

- No agent may enter production where authority cannot be traced to a valid delegator, or where the agent does not hold its own identifiable, least-privilege credentials (identity/credentials checkpoint: Gate 2; re-confirmed at Gate 6).
- No autonomous action may proceed where the required containment mechanism is not demonstrated (containment/kill-switch checkpoint: Gate 2; re-confirmed at Gate 6).
- Any agent with persistent/shared memory must pass the memory-control checkpoint (Gate 2; re-confirmed at Gate 6).
- Any agent invoking external tools must pass the tool/interface authority checkpoint (Gate 2; re-confirmed at Gate 6).
- Any multi-agent system must pass the delegation-chain checkpoint, including recursion, depth, budget and termination controls (Gate 2; re-confirmed at Gate 6).
- Per-action human authorisation is required before any irreversible or high-consequence action, unless an equivalent control is formally authorised at the relevant Council gate (evidenced at Gate 2; re-confirmed at Gate 6) and recorded in AIG-AGT-04. Post-action review and monitoring-only modes are permitted only for non-consequential, reversible actions. "Unsure" about per-action review is treated as "No".
- Consequential actions must have a tested rollback, or a defined compensating action and owner where reversal is not possible (Playbook F.7.13), before go-live (rollback checkpoint: Gate 2; re-confirmed at Gate 6).
- Any model-routing or dynamic model-selection mechanism must be included in the governed configuration.

Lean routing requirements

The Agentic Triage Function is a routing mechanism, not an additional approval layer.

- It is invoked from the existing AI Intake/triage flow when a "can it act?" trigger or equivalent agentic trigger is present.
- It records the minimum information needed to determine agency, authority and control requirements, then writes the result to the Agent Record/ASBOM.
- It must not require users to duplicate information already captured in the AI Register, risk assessment or specialist assurance records.
- The output is a route and required evidence set; approval remains with the existing authorised decision forum.

Annex — Triage contract and record joins

PROPOSED INTEGRATION ONLY — This annex supplements the preserved source artefact; it does not amend or replace original text or tables and is not approved policy or permission to operate.

Identity lifecycle

- A provisional discovery/case reference may support disclosure of AI found in use before registration; it is not AIR-ID. Governance assigns permanent AIR-ID when registering in AI Register (AIG-INV-04). Pair each individual agent's stable AG-ID with AIR-ID; one system may have multiple AG-IDs.
- Record exact AIG-AGT-04 ASBOM version/revision and identifiable snapshot/effective point, not an unspecified "latest".

Separate measures and authority source

- Five dimensions—consequence, autonomy, authority, reach and controllability—form an evidenced profile with capability flags and mandatory tier floors. They do not replace AGPI (AIG-ASS-01) governance priority or Risk Worksheet (AIG-ASS-02) inherent/residual risk tiers; none is approval.
- ASBOM (AIG-AGT-04) is source for capability, permissions, components, controls and delegations. AIG-AGT-05 derives from ASBOM Authority & Delegations edges and helps examine paths, ceilings, revocation and reach; resolve authority against exact ASBOM version. Triage output is a route/evidence set, not a new approval layer; the delegated forum decides.

Full re-entry and action evidence

- Discovered AI already in use enters full Intake (AIG-INV-03) and full triage, never presumed Fast Track. Material change to purpose, model, data, affected people, connections, capability, permissions, authority, autonomy, environment or controls requires full triage re-entry before changed consequential use, refreshing AGPI/risk/specialist review as applicable.
- AIG-AGT-06 records unique Action ID and execution evidence linked to AIR-ID, AG-ID, exact ASBOM version, authority and result. Cross-reference in AIG-OPS-03/AIG-OPS-02/AIG-DEC-04 where applicable. It is not a gate decision or permission grant.

Status of this proposed copy

Proposed map handoff — does not score or route

AIG-INV-05 is an optional proposed context map, not an additional triage score, control, gate or intake. If used, triage may cite a dated UC/CAP/system relationship as a source pointer; independently verify any fact relevant to consequence, autonomy, authority, reach, controllability, multipliers or mandatory floors. Governance issues AIR-ID through AIG-INV-04, plans prospective gates in AIG-DEC-04, and records actual events/conditions separately. The AIG-DEC-01 Governance Gate Map remains forum-routing configuration, not this capability map.

Proposed integrated draft; not approved and not effective. Original baseline content and tables are retained above. Formal approval, owner confirmation and document-control updates remain for the competent Council authority.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 56-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-001 — UK General Data Protection Regulation (UK GDPR), as amended	Where personal data is processed; test each UK GDPR provision separately.	https://www.legislation.gov.uk/eur/2016/679/contents
REQ-032 — Regulation (EU) 2024/1689 (EU AI Act), consolidated and amended by enacted Regulation (EU) 2026/1744	Only after case-specific EU territorial, actor-role and material scope assessment.	https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:02024R1689-20260727 ; amending Regulation (EU) 2026/1744: https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:32026R1744
REQ-040 — ISO/IEC 42001:2023 Artificial intelligence management system	Proposed management-system design mapping; licensed clause review and adoption needed.	https://www.iso.org/standard/42001

NIST-AI-100-1 — NIST AI Risk Management Framework 1.0 (NIST AI 100-1, January 2023)

Voluntary alignment reference; not certification or statutory conformity

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10>