

PROPOSED INTEGRATED DRAFT — NOT APPROVED*For Council review only · Not authorised, effective or operational until formally approved*

AIG-AGT-03 · The Council · AI GOVERNANCE

Agentic Triage


Function specification for the two-function triage route

Field	Controlled value
Document title	Agentic Triage: Function Specification
Status	Working draft; pre-adoption
Author	AI Governance Lead
Date	19 September 2026
Purpose	Specify how the triage route handles agents: one front door, two functions, dynamic assessment
Built to	Agentic AI Governance Design Principles (AIG-AGT-01); realises the recommendation in the Agentic Triage design note
Uses / writes	Agentic Classification Reference (vocabulary); writes Autonomy Level and Agency Tier to the Agent Record / ASBOM
Source	OWASP Top 10 for Agentic Applications 2026 and Agent Control Standard; uploaded agentic triage method
Classification	OFFICIAL

1. What this specifies

The design note recommended keeping one admission route and adding a distinct, dynamic agentic triage for agents. This document specifies how that works: the gate, the two functions, the steps inside the agentic function, the output and how it connects to the rest of the layer. It does not repeat the options analysis; it builds the chosen one.

2. The route at a glance



Action risk is re-evaluated at runtime; a system moves between risk states as its action and authority change.

One front door and one AIR-ID. Every AI system is priority-triaged; anything that can act is additionally agency-triaged; the two readings combine to route it.

Run the system intake once under AIR-ID, then identify each materially distinct use/action with its own UC-ID. Determine priority, risk, route and delegated decision scope per use; do not let a lower-risk system baseline suppress a higher-risk UC. Only an actually action-capable configuration receives the additional agentic/action/authority triage; ordinary tools remain non-agentic while still receiving UC-ID scope.

3. Function one: system and priority triage

This is the existing triage, run for every AI system, unchanged except for one addition.

- Priority and risk: AGPI scores governance priority and the risk worksheet sets inherent and residual risk tiers, exactly as now.
- The gate: one question is added at the top: can it act? If no, the system continues down the standard AI review path and the agentic function does not fire. If yes, it is an agent and function two runs.

4. Function two: the agentic triage

Fires only for agents. It runs in this order, and it governs the authority envelope, what the agent can do, not what it is meant to do.

1 Freeze the system as an object: complete the Agentic System Snapshot (purpose, owner, environment, models, agents, memory, tools, data, interfaces, human involvement, persistence, delegation, authority). Establish what exists before scoring anything.

2 Map the action chain and mark every point where the system can change reality, so governance effort lands where consequences occur.

- 3 Separate cognitive capability from action capability, and record the capability vector. A clever model with little agency is not the concern; a modest model wired to payments and email is.
- 4 Establish the authority boundary: the maximum the agent can do without asking anyone. Look for capability exceeding intended authority.
- 5 Score the five agency dimensions in the Agentic Classification Reference (AIG-AGT-02): consequence, autonomy, authority, reach and controllability. Record a 0–5 profile, not a sum or average. Controllability is reverse-scored: 5 means hardest to stop or reverse. Flag persistence, delegation and the other agency multipliers separately.
- 6 Run the worst plausible chain: starting from a legitimate task, the most consequential chain of actions the architecture permits.
- 7 Test the boundaries (authority, tool, data, identity, delegation, memory, instruction, escalation, persistence, termination), not just the model.
- 8 Run the kill-switch test: not could you stop it, but show me. Demonstrate stopping it, revoking credentials, disabling tools, handling actions in flight, and reconstructing what it did.
- 9 Assign the agency tier T0 to T5 from the profile, and record it.

5. Dynamic, not one-time: the three linked assessments

Assessment join — retain the system AIR-ID and cite the applicable UC-ID in each risk, authority, gate, action and incident record. Shared system evidence may be reused only with an explicit applicability check; use-specific outcome, action scope, residual risk and decision must remain separately traceable.

A single classification at deployment is treated as insufficient, because the same agent moves between risk states. So the agentic function produces three linked assessments rather than one score.

- System risk: what could the overall system cause? Set at the gate and at reauthorisation.
- Action risk: how risky is this particular action? Evaluated per consequential action at runtime.
- Authority risk: how much power has been delegated to it? Read from the Authority Graph and revisited when authority changes.

The worked case: an HR agent is low risk summarising CVs, higher ranking candidates, and high autonomously rejecting them. The agent has not changed; its action and authority have. Action risk is where that is caught, which is why it is a runtime function, not a gate score.

6. Output and routing

The tier routes the agent to a governance pathway. Because the five dimensions form a profile, a high-consequence or hard-to-contain reading can escalate the pathway regardless of the other scores. On the reverse-scored controllability scale, harder to contain means a higher score.

Agency tier	Governance pathway
T0 informational	No further agentic review; standard record only
T1 assisted	Standard AI review (existing pathway)
T2 bounded agent	Agentic controls review
T3 consequential agent	Agentic controls review, higher gate
T4 high-agency	High-agency review
T5 exceptional / high-consequence	Executive and safety escalation

7. How it connects to the rest of the layer

- One record: everything enters the same route and keeps one AIR-ID; the priority reading comes from AGPI on the Register (AIG-INV-04).
- Writes to the ASBOM: the agentic function writes the Autonomy Level and the Agency Tier onto the Agent Record / ASBOM, and its boundary and kill-switch results become evidence there.
- Routing: the multi-board router gains the agentic pathways above; it is extended, not replaced.

8. Deployment-time now, runtime next

Consistent with the design principles and Part 2, the function is built in two parts.

- Buildable now: the gate, the snapshot, the profile and multipliers, the worst-plausible-chain, the boundary and kill-switch checks, the tiering and the routing.
- Production gate: per-action risk evaluation and continuous authorisation are required for consequential autonomy where the approved authority envelope depends on them. Later platform delivery is a design roadmap, not evidence or permission to operate. A bounded, supervised pilot may proceed under its specifically approved conditions; production Approved/Active requires required controls Implemented and Evidenced, or an effective, time-bounded compensating control explicitly accepted by the relevant delegated Council authority and recorded in AIG-DEC-03/AIG-DEC-04. If neither exists, keep the system outside consequential production use.

9. The first-meeting checklist

For the first governance conversation about any agent, these ten questions are answered before moving on. They are the human-readable form of the agentic triage.

- 1 What is the agent's purpose?
- 2 What can it actually do, not what was it intended to do?
- 3 What systems and data can it access?
- 4 What can it change in the real world?
- 5 What authority has been delegated to it?
- 6 Can it act without a human?
- 7 Can it persist, remember, delegate or create additional actors?
- 8 What is the worst plausible chain of actions?
- 9 Can we stop it immediately and prove what it did?

AIG-ASS-05 Who is accountable when it causes a consequential outcome?

10. Status and source

Working draft, pre-adoption; nothing evidenced until operated. The method (snapshot, action chain, five dimensions, multipliers, boundary and kill-switch tests, T0 to T5, three linked assessments, the ten questions) is the framing of the uploaded triage analysis, adopted to satisfy OWASP 2026 controls: identity and privilege (ASI03), tool misuse (ASI02), cascading failures (ASI08), human-agent trust (ASI09) and rogue agents (ASI10), with the dynamic stance from the Agent Control Standard (1 September 2026), verified against genai.owasp.org on 19 September 2026.

Working draft. Two functions, one route, one register; writes to the Agent Record / ASBOM.
Operational decision rules

The Agentic Triage Function must produce a reproducible AIR-ID-linked assessment. The assessor records a 0–5 score, rationale and evidence for each of the same five dimensions: consequence, autonomy, authority, reach and controllability. Persistence, delegation and other agency multipliers are recorded separately as capability flags and control triggers; they do not replace autonomy or authority scores. Apply the highest applicable mandatory tier floor, record the resulting route and required controls, and write the permitted Autonomy Level and Agency Tier to the Agent Record / ASBOM (AIG-AGT-04). Retain the scored profile and evidence reference with the governance case.

Use AIG-ASS-11 to record proportionate ASI01–ASI10 applicability and tests for action-capable agents at any base risk tier, and AIG-AGT-04 for each required runtime control state and evidence. Set AIG-OPS-02 monitoring thresholds before go-live; AIG-AGT-06 records consequential actions. AI Assurance Board views are assurance opinions; the decision is made under the Council delegation.

Minimum control gates

- No agent may enter production where authority cannot be traced to a valid delegator.
- No autonomous action may proceed where the required containment mechanism is not demonstrated.
- Any agent with persistent/shared memory must pass the memory-control gate.
- Any agent invoking external tools must pass the tool/interface authority gate.
- Any multi-agent system must pass the delegation-chain gate, including recursion, depth, budget and termination controls.
- Any agent performing irreversible or high-consequence action must have an explicit human oversight/approval mode or a formally authorised equivalent control.
- Any model-routing or dynamic model-selection mechanism must be included in the governed configuration.

Lean routing requirements

The Agentic Triage Function is a routing mechanism, not an additional approval layer.

- It is invoked from the existing AI Intake/triage flow when a “can it act?” trigger or equivalent agentic trigger is present.
- It records the minimum information needed to determine agency, authority and control requirements, then writes the result to the Agent Record/ASBOM.
- It must not require users to duplicate information already captured in the AI Register, risk assessment or specialist assurance records.
- The output is a route and required evidence set; approval remains with the existing authorised decision forum.

Annex — Triage contract and record joins

PROPOSED INTEGRATION ONLY — This annex supplements the preserved source artefact; it does not amend or replace original text or tables and is not approved policy or permission to operate.

Identity lifecycle

- A provisional discovery/case reference may support disclosure of AI found in use before registration; it is not AIR-ID. Governance assigns permanent AIR-ID when registering in AI Register (AIG-INV-04). Pair each individual agent's stable AG-ID with AIR-ID; one system may have multiple AG-IDs.
- Record exact AIG-AGT-04 ASBOM version/revision and identifiable snapshot/effective point, not an unspecified "latest".

Separate measures and authority source

- Five dimensions—consequence, autonomy, authority, reach and controllability—form an evidenced profile with capability flags and mandatory tier floors. They do not replace AGPI (AIG-ASS-01) governance priority or Risk Worksheet (AIG-ASS-02) inherent/residual risk tiers; none is approval.
- ASBOM (AIG-AGT-04) is source for capability, permissions, components, controls and delegations. AIG-AGT-05 derives from ASBOM Authority & Delegations edges and helps examine paths, ceilings, revocation and reach; resolve authority against exact ASBOM version. Triage output is a route/evidence set, not a new approval layer; the delegated forum decides.

Full re-entry and action evidence

- Discovered AI already in use enters full Intake (AIG-INV-03) and full triage, never presumed Fast Track. Material change to purpose, model, data, affected people, connections, capability, permissions, authority, autonomy, environment or controls requires full triage re-entry before changed consequential use, refreshing AGPI/risk/specialist review as applicable.
- AIG-AGT-06 records unique Action ID and execution evidence linked to AIR-ID, AG-ID, exact ASBOM version, authority and result. Cross-reference in AIG-OPS-03/AIG-OPS-02/AIG-DEC-04 where applicable. It is not a gate decision or permission grant.

Status of this proposed copy

Proposed map handoff — does not score or route

AIG-INV-05 is an optional proposed context map, not an additional triage score, control, gate or intake. If used, triage may cite a dated UC/CAP/system relationship as a source pointer; independently verify any fact relevant to consequence, autonomy, authority, reach, controllability, multipliers or mandatory floors. Governance issues AIR-ID through AIG-INV-04, plans prospective gates in AIG-DEC-04, and records actual events/conditions separately. The AIG-DEC-01 Governance Gate Map remains forum-routing configuration, not this capability map.

Proposed integrated draft; not approved and not effective. Original baseline content and tables are retained above. Formal approval, owner confirmation and document-control updates remain for the competent Council authority.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 51-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-001 — UK General Data Protection Regulation (UK GDPR), as amended	Where personal data is processed; test each UK GDPR provision separately.	https://www.legislation.gov.uk/eur/2016/679/contents
REQ-032 — Regulation (EU) 2024/1689 (EU AI Act), consolidated and amended by enacted Regulation (EU) 2026/1744	Only after case-specific EU territorial, actor-role and material scope assessment.	https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:02024R1689-20260727 ; amending Regulation (EU) 2026/1744: https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:32026R1744
REQ-040 — ISO/IEC 42001:2023 Artificial intelligence management system	Proposed management-system design mapping; licensed clause review and adoption needed.	https://www.iso.org/standard/42001

NIST-AI-100-1 — NIST AI Risk Management Framework 1.0 (NIST AI 100-1, January 2023)

Voluntary alignment reference; not certification or statutory conformity

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10>