

AIG-AIMS-11 — The Council — AI System Impact Assessment: Coverage Mapping

AI System Impact Assessment — Coverage Mapping

AI Governance Playbook (Level 5 document) · ISO/IEC 42001:2023 clause 6.1.4 and Annex A controls A.5.2 to A.5.5 · v1.4 draft

This note shows how the Council satisfies the clause 6.1.4 requirement to assess the impacts of AI systems on individuals, on groups of individuals, and on society. The requirement is met by the Council's existing impact assessments acting together; this mapping states which assessment covers which dimension, so the coverage is explicit and auditable. Draft for Council review; clause references to be validated against the Council's licensed copy of the standard.

How the requirement is met

Clause 6.1.4 requires an assessment of the potential consequences of AI for individuals, groups of individuals and society. The Council meets this through the impact assessments completed within the governance pathway, applied to the risk tier and context of each system. Together they provide the three required lenses. Environment and security are added as supporting lenses (Proposed — for Council confirmation):

Impact dimension	Assessment(s) that cover it	What is assessed
Individuals	DPIA (AIG-ASS-05) and Human Rights Assessment (AIG-ASS-07)	Data-protection risks to individuals and the automated-decision safeguards; engagement of Convention rights and whether interference is lawful, necessary and proportionate.
Groups of individuals	Equality Impact Assessment (AIG-ASS-06)	Differential impact and differential accuracy across people sharing protected characteristics under the Equality Act and the Public Sector Equality Duty.
Society	Responsible AI Assessment (AIG-ASS-04)	Public value and proportionality of the use, and wider societal effects, through the public-benefit and accountability themes.
Environment	Responsible AI Assessment (AIG-ASS-04, theme 2) and Supplier AI Due-Diligence Questionnaire (AIG-ASS-08, question 4.4)	Whether the energy and compute use is proportionate (could a smaller or simpler model do the job), the supplier's energy and carbon information, data-centre locations and hosting region, and supplier sustainability commitments.
Security	AI Security Review Checklist (AIG-ASS-11), including the Section 8 threat model	Threats to the system and to the people who rely on it: assets, entry points, likely misuse and AI-specific threats, with controls and residual concerns.

Process and documentation (A.5.2 to A.5.5)

- **Coverage key** — report AIR-ID (system) and UC-ID (materially distinct use) for every impact/assurance coverage conclusion. Show the affected population, assessed scope, evidence reference and current per-use risk/decision position; mark unknown or conflicting links unverified. This coverage map is not itself an assessment or approval.
- Which assessments are required, and their depth, is set at intake by the Governance Checklist (AIG-ASS-03) against the applicable use-specific risk tier. Risk tier sets the depth of assessment; equality screening (PSED due regard) and DPIA screening apply to every in-scope use regardless of tier. Trigger and scope impact assessment for each materially distinct in-scope use (UC-ID) beneath its system AIR-ID.
- Record the system identity by AIR-ID and each assessment conclusion against the relevant UC-ID, with references to AIG-ASS-02/source assessment and the delegated AIG-DEC-03 decision or approved minutes. The system Register row is not use approval.
- Assessments are repeated on material change or re-triage, so the impact view stays current across the lifecycle.

Known limitation

The societal dimension, which ISO/IEC 42001 Annex A control A.5.5 (assessing societal impacts of AI systems) addresses, is the lightest of the three: the Responsible AI Assessment addresses public benefit and proportionality but not a full societal-impact analysis. Where a system's societal effects are significant (for example large-scale, novel, or affecting vulnerable groups or public trust), the Responsible AI Assessment should be extended with an explicit societal-impact prompt, or a dedicated societal-impact assessment added. This is recorded as an improvement action against control A.5.5 (AIG-AIMS-03 IMP-17) and in the Statement of Applicability (AIG-AIMS-02).

Document control: AIG-AIMS-11 | Version: 1.4 draft | Owner: AI Governance Lead | Status: Draft for Council review

Change history: 1.1 draft, 29 September 2026: QA remediation — applied findings QA-074, QA-077, QA-086, QA-109 (see QA change log).

1.2 draft, 30 September 2026: v3.9.1: standards cross-reference note updated to the 54-entry requirement inventory (AIG-AIMS-05 REQ-054, Civil Contingencies Act 2004, added).

1.3 draft, 1 October 2026: v3.9.3: standards gap closure against NCSC, NIST AI 600-1 and ISO/IEC 42005 (see QA change log, sheet “v3.9.3 changes”): Environment and Security lens rows added to the coverage table (C4); new annex “Risk domain crosswalk” mapping the fourteen Playbook §4.3 risk domains to NIST AI 600-1 risks and ISO/IEC 42005 impact dimensions, with the NIST risks that have no domain home (C8); standards cross-reference adds REQ-055 and cites the 55-entry requirement inventory. Proposed — for Council confirmation.

1.4 draft, 1 October 2026: v3.9.4: the risk domain crosswalk note now points to the open check of the ISO/IEC 42005 mapping against the Council's licensed copy (AIG-AIMS-05 REQ-042) (see QA change log, sheet “v3.9.4 changes”, D3).

PROPOSED INTEGRATION ADDITIONS — REVIEW REQUIRED

Proposed map boundary in impact coverage

AIG-INV-05 can be a discovery aid for use-case/capability/system relationships, but it does not replace AIG-AIMS-11 impact-assessment coverage or specialist analysis. A map relationship is not evidence that a group is affected or unaffected. Cite the underlying assessment sources, retain uncertain/unknown relationships for review and use current AIG-INV-04 identity/current state plus relevant actual AIG-DEC-04 Events only as context.

NOT APPROVED. Additive review annex only; original paragraphs and tables above are retained. These proposals do not record live facts, completed actions, delegated authority, Council adoption or approval. Until an approved migration, the current Register and Gate Log practice remains authoritative; the AIG-INV-04/AIG-DEC-04 joins below are proposed design only.

Coverage remains design-only

This mapping documents the proposed complementary lenses—DPIA/HRA for individual impacts, EIA for group effects, and Responsible AI assessment for societal effects—subject to the source document's stated limitation on fuller societal analysis. It does not establish that an assessment was triggered, completed, approved, comprehensive or effective for any system.

AIR-ID evidence join

For each relevant system, use proposed AIG-INV-04 AIR-ID to relate the applicability/assessment references in the Assessment summary worksheet to the actual AIG-ASS-04/AIG-ASS-05/AIG-ASS-06/AIG-ASS-07 assessment records and an Evidence ID/version/location in the AIG-INV-04 Evidence Index. Keep assessment content in its controlled source, with specialist approvals/limitations retained there; a reference, checkbox, or template title alone is not completion evidence.

AIG-DEC-04 gate / reassessment trail

If an assessment is a prerequisite or action of a real gate, relate its planned review to AIG-DEC-04 Gate Plan, the actual dated approval/reassessment event to Gate Events, and each required follow-up to Gate Conditions. Evidence and authorised decision remain in source assessment records and AIG-DEC-03/native minutes. A planned gate, blank condition, or gate status does not demonstrate impact assessment or specialist sign-off.

Change and coverage gap

Material purpose, model/data, population, context, autonomy, service, supplier, legal or impact change should trigger documented reassessment according to approved process. Retain the AIG-INV-04 change record, changed assessment version and actual gate event/condition references if applicable; update the crosswalk only after a justified review. For significant societal effects, retain explicit supplemental analysis and evidence rather than claiming the Responsible AI template alone covers all societal consequences.

Use in assurance

Report assessment coverage using a defined population, applicability rationale, as-at date and independent evidence review. Link any issue to AIG-AIMS-03 risk/treatment or AIG-AIMS-08 where an actual nonconformity is established; keep complaints/incidents in their source record. Do not infer population-level coverage from the existence of four templates.

Validation before claims

Validate the interpretation of clause 6.1.4 against the licensed standard, assessment triggers and specialist sign-off rules with DPO, Equality, Human Rights, Legal and relevant services. Confirm the proposed AIG-INV-04/AIG-DEC-04 join and protected-data/access controls before adoption.

Annex: Risk domain crosswalk (Proposed — for Council confirmation)

This table maps each of the fourteen risk domains in Playbook §4.3 to the related risks in NIST AI 600-1, the Generative AI Profile of the NIST AI Risk Management Framework (July 2024; AIG-AIMS-05 REQ-055), and to impact dimensions used in ISO/IEC 42005:2025 AI system impact assessment (REQ-042). Use it as a completeness prompt when assessing a generative AI use; it is a design mapping, not evidence that a risk has been assessed.

Playbook §4.3 risk domain	Related NIST AI 600-1 risks	ISO/IEC 42005 impact dimensions
4.3.1 Bias and Fairness	Harmful bias or homogenisation	Fairness

Playbook §4.3 risk domain	Related NIST AI 600-1 risks	ISO/IEC 42005 impact dimensions
4.3.2 Transparency and Explainability	Human-AI configuration; information integrity	Explainability and transparency
4.3.3 Privacy and Data Protection	Data privacy	Privacy
4.3.4 Security	Information security; value chain and component integration	Safety; reliability
4.3.5 Data Quality and Governance	Information integrity; harmful bias or homogenisation; value chain and component integration	Reliability; fairness
4.3.6 Model Performance and Drift	Confabulation; information integrity	Reliability
4.3.7 Hallucination and Content Accuracy	Confabulation; information integrity; dangerous, violent or hateful content	Reliability; safety
4.3.8 Automation Bias and Over-Reliance	Human-AI configuration; confabulation	Accountability; safety
4.3.9 Accountability and Governance	Human-AI configuration; value chain and component integration	Accountability
4.3.10 Supplier and Third-Party	Value chain and component integration; intellectual property	Accountability; explainability and transparency
4.3.11 Legal and Regulatory Compliance	Data privacy; intellectual property; harmful bias or homogenisation	Accountability; privacy; fairness
4.3.12 Reputational and Public Trust	Information integrity; dangerous, violent or hateful content	Accountability; explainability and transparency
4.3.13 Operational and Service Delivery	Value chain and component integration; human-AI configuration	Reliability; safety
4.3.14 Ethical and Public Value	Harmful bias or homogenisation; human-AI configuration; information integrity	Accountability; fairness

NIST AI 600-1 risks with no domain home, and where they are handled: chemical, biological, radiological or nuclear (CBRN) information or capabilities, and obscene, degrading or abusive content, are prohibited uses under the acceptable use policy (Playbook §5.4.3); environmental impacts are covered by the Responsible AI Assessment (AIG-ASS-04, theme 2) and the supplier questionnaire (AIG-ASS-08, question 4.4), and by the Environment lens above.

The ISO/IEC 42005 mapping is to be validated against the Council's licensed copy of the standard before it is relied on (licensed-copy check pending: AIG-AIMS-05 REQ-042). The NIST mapping is reference only; NIST AI 600-1 is voluntary US guidance, not a statutory requirement.

Standards and legal applicability — proposed QA cross-reference
PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 55-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-005 — Equality Act 2010, including Public Sector Equality Duty section 149	Public authority due regard under s.149; case-level equality screening is not risk-tier gated.	https://www.legislation.gov.uk/ukpga/2010/15/section/149
REQ-006 — Human Rights Act 1998	Public-authority compatibility under s.6; assess Convention	https://www.legislation.gov.uk/ukpga/1998/42/contents

REQ-040 — ISO/IEC 42001:2023 Artificial intelligence management system	rights engaged by the case. Proposed management-system design mapping; licensed clause review and adoption needed.	https://www.iso.org/standard/42001
REQ-042 — ISO/IEC 42005:2025 AI system impact assessment	Guidance for AI impact-assessment design; validate edition and selected use.	https://www.iso.org/standard/42005
REQ-055 — NIST AI 600-1 Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (July 2024)	Voluntary guidance (US); generative AI risk prompts and actions; used in the risk domain crosswalk annex. Reference only, not statutory.	https://doi.org/10.6028/NIST.AI.600-1