

AI Governance Checklist

The Council · AI Governance Playbook (Level 5 document, §3.8 / §3.10.1)

The assessments, reviews and approvals required for an AI system, by risk tier.

Determined by the AI Governance Lead at intake Step 4 (governance pathway assignment) and issued to the Service Owner with the Step 5 intake acknowledgement. It sets out the assessments, reviews and approvals required for the assigned risk tier before the system may be deployed. The assigned tier sets the depth of assurance and the level of AI Assurance Board involvement; the decision is made by the officer or Council forum holding confirmed delegation, and the AI Assurance Board provides assurance, challenge and recommendations rather than approval (Playbook §3.8, §3.10.1; Gate Map AIG-DEC-01). Where classification is uncertain, apply the higher tier. It is primarily an internal routing and assurance checklist: the Service Owner is asked only for items marked as their responsibility. Existing evidence is referenced by AIR-ID rather than recreated or attached again.

1. System and Assigned Tier

System / model name	
AIR-ID	AIR-XXXX
UC-ID / purpose / outcome	One materially distinct use per assessment unit; repeat as needed.
Service Owner	
Assigned risk tier	Low / Medium / High / Critical
Delegated decision authority	Officer or Council forum holding confirmed delegation (see table below)

Assessment unit — complete this checklist for the named UC-ID and its specific purpose, population and decision/action. Set use-specific priority, risk tier, controls and authority route; repeat for materially different uses under the same AIR-ID. Checklist completion, tier or recommendation is not the delegated decision.

2. Approval and Assurance by Tier

Tier	Delegated authority deciding each UC-ID	Required governance activities
Low	The officer or forum with confirmed delegation (Service Owner/Service Manager only if delegated)	Register entry, basic assessment, policy compliance review
Medium	Competent delegated forum/officer under confirmed delegation (AI Governance Working Group detailed review; AI Assurance Board reviews and provides an assurance opinion/recommendation)	Structured risk assessment, Responsible AI Assessment, governance review, monitoring requirements; for resident-facing generative AI, an adversarial test (AIG-ASS-11 Section 7) before go-live (Proposed — for Council confirmation)
High	Competent delegated forum/officer under confirmed delegation (Assurance Board formal recommendation and single assurance opinion)	Enhanced assessment; DPIA where statutory/Council screening requires; EIA; Legal and Security review; Responsible AI Assessment; Human Rights Assessment where Convention rights may be engaged; independent assurance review (Playbook §4.5.3, §4.5.9)

Tier	Delegated authority deciding each UC-ID	Required governance activities
Critical	Competent delegated forum/officer under confirmed delegation, with Senior Responsible Owner (SRO) and Executive Sponsor oversight; escalation to the relevant executive member or committee (Cabinet Member/Cabinet or the equivalent committee) and the Audit and Risk Committee (or equivalent)	All High controls plus mandatory Human Rights Assessment, executive oversight, independent challenge and independent assurance (Playbook §3.10.2) and formal reporting

At every tier, screen the UC-ID for possible equality and Convention-rights engagement and record the reasoned outcome (AIG-ASS-06, AIG-ASS-07); refer possible or uncertain Convention-rights engagement to Legal Services. Where a current assessment already covers this kind of use, the screening may be done by reference: cite its Evidence ID, confirm the use is within its scope and record the reasoned outcome (Playbook §4.6; for Light-touch, the AIG-INV-02 Part C box). The AI Assurance Board provides assurance, challenge and recommendations; it does not approve or decide unless formally delegated.

3. Required Items — Tick and Track

The AI Governance Lead determines the applicable items and records N/A where justified. The Service Owner supplies only the evidence assigned to them. Capture evidence once and reference the source by AIR-ID; do not duplicate evidence across the checklist, Register, Gate Log or Decision Record.

Governance item	Required?	Owner	Complete?	Date
AI Register entry				
Responsible AI Assessment (AI Governance Lead QA)				
DPIA (DPO advice recorded; Controller/SRO decision)				
Equality Impact Assessment (Equalities QA)				
Human Rights Assessment (Legal Services advice): screening at every tier; full assessment for Critical, and for High where Convention rights may be engaged				
AI Security Review Checklist (Security assurance)				
Independent assurance review: High (Playbook §4.5.3, §4.5.9); Critical, with independent challenge (§3.10.2)				
Threat model completed (AIG-ASS-11 Section 8): High and Critical, and any action-capable use (Proposed — for Council confirmation)				

Governance item	Required?	Owner	Complete?	Date
Adversarial test for resident-facing generative AI at Medium (AIG-ASS-11 Section 7), completed before go-live (Gate 6) (Proposed — for Council confirmation)				
Legal advice recorded				
Model Card				
ATRS reportability checked				
Monitoring plan in place, with agreed thresholds and who agreed them recorded (AIG-OPS-01 section 7; AIG-OPS-02)				
Governance Decision Record (delegated decision, AIG-DEC-03)				
Agentic security and runtime evidence (AIG-ASS-11, AIG-AGT-04), action trace (AIG-AGT-06), monitoring thresholds (AIG-OPS-02) — whenever the AI can act; proportionate to authority and consequence		Information and Cyber Security / Service Owner		

4. Confirmation and Governance Decision

The Service Owner confirms the required evidence is complete. Completion is not approval: deployment of a UC-ID requires its recorded Governance Decision (AIG-DEC-03) by the officer or Council forum holding confirmed delegation, used within all recorded conditions (Playbook Appendix F.7.8). The AI Register Governance Approval Status / Operational Status (system baseline; not use permission) is then reconciled, but a system-level Approved/Active state does not itself permit any use.

Service Owner name	
Signature	
Date	
AI Governance Lead QA name	
QA review recorded / signature	
Date	
Decision-maker (officer or forum with confirmed delegation)	
Governance Decision Record reference	
Delegated decision and date (as recorded in AIG-DEC-03)	

For an action-capable system, apply the AIG-ASS-11 agentic ASI01–ASI10 review at any risk tier. Record applicability, tests and evidence in AIG-ASS-11/AIG-AGT-04, monitoring plan in AIG-OPS-02 and gate outcome in AIG-DEC-03/AIG-DEC-04. The agency tier sets a minimum risk-tier pathway and the higher of that and the risk-tier pathway governs (Gate Map AIG-DEC-01, Agentic pathway); Gate 2 is mandatory for every action-capable use, with the agentic control checkpoints evidenced there and re-confirmed at Gate 6. A light risk classification does not waive an agentic control triggered by actual authority. The AI Assurance

Board provides assurance, challenge and recommendations; the officer or Council forum holding confirmed delegation makes the decision.

Document control

Administrative control record. This artefact remains a working draft until formally approved.

Field	Controlled value
Document ID	AIG-ASS-03
Version	1.10 draft
Status	Draft for Council review
Document owner	AI Governance Lead
Approving body	Pending confirmation — Competent Council authority under Constitution and Scheme of Delegation; AI Assurance Board provides assurance where applicable; exact delegation and publication authority to be verified by Monitoring Officer.
Approval date	Not yet approved
Effective date	Not yet effective
Next review date	To be set on approval: 12 months after the effective date, or sooner following material change
Classification	OFFICIAL

Change history

Version	Date	Change summary	Owner
1.0 draft	6 August 2026	Document-control baseline established following full-suite quality assurance; no substantive content change.	AI Governance Lead
1.1 draft	20 September 2026	Lean checklist rule added: governance-led applicability, Service Owner-only evidence requests, N/A handling and capture-once/reference-by-AIR-ID.	AI Governance Lead
1.2 draft	25 September 2026	Added action-capable security and runtime evidence checks at all baseline risk tiers.	AI Governance Lead
1.3 draft	29 September 2026	QA remediation — applied findings QA-007, QA-058, QA-062, QA-172, QA-175, QA-185, QA-204, QA-221, QA-032, QA-043, QA-130 (see QA change log). v3.7 alignment: agency-tier minimum pathway and Gate 2/Gate 6 agentic control checkpoints referenced for action-capable systems.	AI Governance Lead
1.4 draft	30 September 2026	v3.9: adversarial test for resident-facing generative AI at Medium (AIG-ASS-11 Section 7) added to the Medium tier activities and as a checklist item, completed before go-live (Proposed — for Council confirmation).	AI Governance Lead
1.5 draft	30 September 2026	v3.9.1: standards cross-reference note updated to the 54-entry requirement inventory (AIG-AIMS-05 REQ-054, Civil Contingencies Act 2004, added).	AI Governance Lead
1.6 draft	30 September 2026	v3.9.2 (pilot-readiness fixes from the scenario test; see QA change log): tier table adds independent assurance review at High (Playbook §4.5.3, §4.5.9) and independent challenge and independent assurance at Critical (§3.10.2); checklist row added for the independent assurance review (W-11).	AI Governance Lead
1.7 draft	1 October 2026	v3.9.3: standards gap closure against NCSC, NIST AI 600-1 and ISO/IEC 42005 (see QA change log, sheet “v3.9.3 changes”): section 3 adds the item	AI Governance Lead

Version	Date	Change summary	Owner
1.8 draft	2 October 2026	“Threat model completed (AIG-ASS-11 Section 8)” for High and Critical and any action-capable use (C3); standards cross-reference note cites the 55-entry requirement inventory (AIG-AIMS-05 REQ-055, NIST AI 600-1 Generative AI Profile, added). Proposed — for Council confirmation. v3.9.5: standards cross-reference note cites the 56-entry requirement inventory (AIG-AIMS-05 REQ-056, SI 2026/425 ICO statutory code of practice on AI and automated decision-making, added). Proposed — for Council confirmation.	AI Governance Lead
1.9 draft	3 October 2026	v3.9.6: the pre-go-live “Monitoring plan in place” item now also asks for agreed thresholds and who agreed them to be recorded, matching AIG-OPS-01 v1.10 section 7. Proposed — for Council confirmation.	AI Governance Lead
1.10 draft	3 October 2026	v3.9.7: screening may be done by reference to a current covering assessment, with its Evidence ID, a scope check and a reasoned outcome (Playbook §4.6). Proposed — for Council confirmation.	AI Governance Lead

PROPOSED INTEGRATION ADDITIONS — REVIEW REQUIRED

Proposed map use — context only

Use the proposed AIG-INV-05 map only to orient the reviewer to intended service outcomes, capabilities and documented system/component relationships. It is neither an assessment nor evidence of control effectiveness, legal applicability, risk acceptance or approval. Checklist completion and any map link do not alter the route; consult current state in AIG-INV-04, prospective requirements in the AIG-DEC-04 Plan, actual dated outcomes in AIG-DEC-04 Events and event-linked Conditions, and authoritative decisions in AIG-DEC-03/minutes.

This integrated draft includes revisions to body content above; the annex below is supplementary proposed content for Council review. This draft and annex do not record live facts, completed actions, delegated authority, Council adoption or approval.

PROPOSED identity, state and evidence-reference rule

Use the permanent AIR-ID issued in AIG-INV-04. Read canonical system identity and current governance/operational state from AIG-INV-04; do not re-key or independently maintain a competing status. If not yet registered, show only the provisional intake reference and leave AIR-ID/state pending.

For each applicable requirement, record a distinct condition/reference, one accountable owner, due date, required evidence source/reference, evidence date/version as applicable, verification outcome and verifier/date. A general completion tick is insufficient to close an approval condition. Evidence is referenced, not duplicated.

The Service Owner confirms assigned evidence; the Governance Lead determines applicability and records reasoned N/A. Completion is not authorisation. The delegated decision authority records its decision in AIG-DEC-03 and cites the authoritative native meeting minutes or controlled written decision; AIG-DEC-04 carries dated planned and actual gate events.

PROPOSED condition closure and deployment gate

Carry each condition from AIG-DEC-03 into AIG-DEC-04 with owner, due date, evidence and closure verifier. Record actual closure date separately from the planned date. An unverified or undocumented condition remains open.

Before go-live reconcile every required condition. A condition may remain open only if expressly accepted by the authorised decision-maker within confirmed delegation, with owner, due date, residual-risk treatment and escalation/review route. No checklist tick waives statutory/legal duties, mandatory controls or required specialist decisions.

Review points — not yet decided

Confirm condition identifiers and whether AIG-DEC-04 supports per-condition evidence and verifier fields; confirm delegated limits and authoritative minutes/written-decision repository.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 56-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-040 — ISO/IEC 42001:2023 Artificial intelligence management system	Proposed management-system design mapping; licensed clause review and adoption needed.	https://www.iso.org/standard/42001
REQ-041 — ISO/IEC 23894:2023 AI — Guidance on risk management	Voluntary risk-management guidance; proposed alignment, not statutory conformity.	https://www.iso.org/standard/77304.html