

Responsible AI Assessment

The Council · AI Governance Playbook (Level 5 document, §4.6.4)

A structured self-assessment of an AI system against The Council's eight Responsible AI themes.

A self-assessment completed by the Service Owner for all Medium, High and Critical risk AI systems. The AI Governance Lead conducts a quality assurance review before it is finalised. Recommended at intake for Medium risk systems; required in the pre-approval package for High and Critical systems (Playbook §4.6.4).

PROPOSED — legal review only; draft clarification, not an adopted or live Council policy change.

This Responsible AI Assessment and its Medium/High/Critical completion rules are internal governance controls; neither the RAIA nor the AGPI tier decides whether a legal duty applies. Record case-level equality and Convention-rights screening where relevant for every AI use irrespective of AGPI tier, using the native EIA/HRA records and their reasoned outcomes; apply the statutory tests independently. Equality Act 2010 section 149: <https://www.legislation.gov.uk/ukpga/2010/15/section/149> ; Human Rights Act 1998 section 6: <https://www.legislation.gov.uk/ukpga/1998/42/section/6>

1. System and Submission Details

System identity	System identity (name, owner, service area): held once in the AI Register (AIG-INV-04) under this AIR-ID; do not re-key here.
UC-ID / purpose / outcome	One materially distinct use per assessment; AGPI priority ____; risk tier ____.
AIR-ID	AIR-XXXX
Risk tier	Medium / High / Critical
Date started	

Use-specific conclusion — assess the named UC-ID, outcome, affected groups and human role. Share system evidence only with an applicability check; benefit, harm, safeguards and residual risk are use-specific. This assessment is not approval.

2. Purpose and Public Benefit

What is the public value? Is the use proportionate to the benefit?

Is the energy and compute use proportionate to the benefit? Would a smaller or simpler model be sufficient?

What energy or carbon information does the supplier publish, and in which region is the system hosted?

(Proposed — for Council confirmation)

Self-rating	Strong / Adequate / Needs work / Not addressed
-------------	--

3. Data Governance and Quality

Are data sources accurate, relevant, lawful and well-managed?

Self-rating

Strong / Adequate / Needs work / Not addressed

4. Transparency and Explainability

Can decisions be explained to residents and staff? Is use disclosed?

Are users told at the start that they are dealing with AI and how to reach a person? Is content produced wholly or substantially by AI checked by a person before release and labelled as AI-assisted where a reader would otherwise assume a person wrote it, and are AI-generated images, audio or video showing real people, places or events clearly labelled and approved before publication (Playbook §5.4.4)? (Proposed — for Council confirmation)

Self-rating

Strong / Adequate / Needs work / Not addressed

5. Fairness and Equality

How are bias and differential impact identified and managed?

Self-rating

Strong / Adequate / Needs work / Not addressed

6. Human Oversight

Who reviews outputs, and can a person override or halt the system?

If a user discloses a risk of harm (self-harm, abuse, crisis or a safeguarding concern), how does the system respond, and how is the user handed over to a person and signposted to help? (Proposed — for Council confirmation)

--

Self-rating

Strong / Adequate / Needs work / Not addressed

7. Accountability and Governance

Who is accountable? How does it sit within governance structures?

How could the system reasonably be misused or used outside its intended purpose, and what prevents or detects this? (Proposed — for Council confirmation)

--

Self-rating

Strong / Adequate / Needs work / Not addressed

8. Privacy and Security

How are personal data and the system protected?

--

Self-rating

Strong / Adequate / Needs work / Not addressed

9. Supplier Governance

If procured, how is the supplier held to the Council's standards?

--

Self-rating

Strong / Adequate / Needs work / Not addressed

10. Overall Outcome and QA

Overall Responsible AI rating	Strong / Adequate / Needs work / Not addressed. Proposed rule (for Council confirmation): if any theme is rated Not addressed, the overall rating is no higher than Needs work.
Key actions before approval	
Evidence references by theme	For each of the eight themes, cite the controlled evidence location, version/date and relevant section.
Remediation action owner(s)	Name the accountable owner for every gap or condition.
Target date(s)	Record a due date for every remediation action.
Closure evidence and status	Open / Closed-verified / Accepted-open / Waived / Unknown (controlled states shared with AIG-DEC-03 and AIG-DEC-04); cite the evidence that demonstrates closure or acceptance.
Escalation trigger and route	Any Needs work / Not addressed rating, overdue action or missing evidence goes to the AI Governance Lead. Unresolved High or Critical matters go to the AI Assurance Board; approval must not be recorded until disposition.

The AI Governance Lead confirms this assessment has been quality assured and is sufficient to support the governance decision.

AI Governance Lead name	
AI Governance Lead signature	
Date	
Service Owner name	
Service Owner signature	
Date	

Document control

Administrative control record. This artefact remains a working draft until formally approved.

Field	Controlled value
Document ID	AIG-ASS-04
Version	1.4 draft
Status	Draft for Council review
Document owner	AI Governance Lead
Approving body	Pending — competent Council authority under confirmed Constitution / Scheme of Delegation; AI Assurance Board provides assurance and is not presumed to hold approval authority.
Approval date	Not yet approved
Effective date	Not yet effective
Next review date	To be set on approval: 12 months after the effective date, or sooner following material change

Field	Controlled value
Classification	OFFICIAL

Change history

Version	Date	Change summary	Owner
1.0 draft	6 August 2026	Document-control baseline established following full-suite quality assurance; no substantive content change.	AI Governance Lead
1.1 draft	6 August 2026	QA revision: added evidence, remediation ownership, due-date, closure and escalation controls.	AI Governance Lead
1.2 draft	29 September 2026	QA remediation — applied findings QA-185, QA-199, QA-200, QA-204, QA-221 (see QA change log). Final tidy-up: self-referencing cross-reference removed from the standards row(s) REQ-033, REQ-034.	AI Governance Lead
1.3 draft	30 September 2026	v3.9.1: standards cross-reference note updated to the 54-entry requirement inventory (AIG-AIMS-05 REQ-054, Civil Contingencies Act 2004, added).	AI Governance Lead
1.4 draft	1 October 2026	v3.9.3: standards gap closure against NCSC, NIST AI 600-1 and ISO/IEC 42005 (see QA change log, sheet “v3.9.3 changes”): new questions under theme 2 (proportionate energy and compute use; C4), theme 4 (telling users they are dealing with AI and labelling AI-generated content; C2), theme 6 (response to a disclosure of risk of harm, hand-over to a person and signposting; C1) and theme 7 (foreseeable misuse; C7); standards cross-reference note cites the 55-entry requirement inventory (AIG-AIMS-05 REQ-055, NIST AI 600-1 Generative AI Profile, added). Proposed — for Council confirmation.	AI Governance Lead

NOT APPROVED — PROPOSED INTEGRATION ANNEX

Integration annex — Responsible AI Assessment

DRAFT FOR COUNCIL REVIEW • 25 September 2026 • Not approved for operational use

Purpose and record boundary

This assessment remains the native Responsible AI assessment and retains its own eight-theme analysis, ratings, remediation and QA record. Do not copy the assessment into the Register: after registration, enter its AIR-ID in this document and keep only the current assurance summary and versioned pointer in AIG-INV-04.

Evidence and governance joins

For each theme, cite the proposed AIG-INV-04 Evidence Index using the specific Evidence ID, source/version, verification or observed date, and controlled location; distinguish missing evidence from a negative finding. An AIG-DEC-04 reference may identify only an actual dated Gate Event or dated Condition (ID and date), not a planned gate or an undated status. The assessment and this annex do not make the governance decision: cite the authorised AIG-DEC-03 decision record or authoritative forum minutes; AIG-DEC-04 may point to that record but cannot approve.

Reassessment and affected people

Revisit the relevant themes when intended use, data, supplier, model/version, configuration, or affected groups materially change. Record current-state changes in AIG-INV-04 and monitoring/reassessment actions in AIG-OPS-02; cite AIG-DEC-04 only for the corresponding dated event/condition. For individual rights and challenge routes, cross-reference AIG-OPS-04 and the relevant DPIA/EIA/HRA without treating a pointer as proof that safeguards operate.

Fundamental-rights coverage: for how this assessment contributes to EU AI Act Article 27 fundamental-rights impact assessment coverage, see the Fundamental Rights Coverage Crosswalk (AIG-ASS-12). The crosswalk is a design mapping, not case evidence.

Agent-enabled use and status

Where consequential delegated action is possible, link the relevant AIR-ID and exact ASBOM/AG-ID version in AIG-AGT-04; use AIG-AGT-05 as a derived authority-graph view and AIG-AGT-06 for action-specific decision records. This does not replace the assessment or grant agent permissions. Draft integration only: no rating, evidence result, legal applicability, approval, or publication is established here.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 55-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-001 — UK General Data Protection Regulation (UK GDPR), as amended	Where personal data is processed; test each UK GDPR provision separately.	https://www.legislation.gov.uk/eur/2016/679/contents
REQ-003 — Data (Use and Access) Act 2025 (DUAA)	Where DUAA-amended UK GDPR/DPA provisions are engaged, including ADM safeguards.	https://ico.org.uk/about-the-ico/what-we-do/legislation-we-cover/data-use-and-access-act-2025/
REQ-005 — Equality Act 2010, including Public Sector Equality Duty section 149	Public authority due regard under s.149; case-level equality screening is not risk-tier gated.	https://www.legislation.gov.uk/ukpga/2010/15/section/149
REQ-006 — Human Rights Act 1998	Public-authority compatibility under s.6; assess Convention rights engaged by the case.	https://www.legislation.gov.uk/ukpga/1998/42/contents
REQ-033 — ICO Guidance on AI and Data Protection and AI Risk Toolkit	Screen the regulatory guidance requirement against this artefact's actual use and record evidence/version.	https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/guidance-on-ai-and-data-protection/
REQ-034 — ICO Explaining Decisions Made with AI guidance	Screen the regulatory guidance requirement against this artefact's actual use and record evidence/version.	https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/explaining-decisions-made-with-artificial-intelligence/
REQ-042 — ISO/IEC 42005:2025	Guidance for AI impact-	https://www.iso.org/standard/

AI system impact assessment	assessment design; validate edition and selected use.	42005
-----------------------------	--	-------