

AI Model Card

The Council · AI Governance Playbook (Appendix A.4)

Complete one card per AI model. Keep it current through the model's operational life and retain it with the AI Register entry.

1. Model Details

Model name	
Version	
AIR-ID	
UC-ID(s) / model use	Purpose/population/decision; AGPI priority ____; risk tier ____.
AI capability (Generative / Agentic / Predictive / Computer Vision / NLP / other ML)	
Lifecycle stage (Pilot / Testing / Live / Retired)	
Responsible owner (Service Owner)	
Supplier / developer	
Date completed	
Next review date	

Use applicability — identify the UC-ID(s) using this model and the intended outcome, population and decision/action. Model-wide facts may be shared; suitability, limitations, risk and mitigations remain use-specific. Do not infer permission for an unlisted or materially changed use.

2. Intended Use

Intended purpose	
Intended users	
In-scope uses	
Out-of-scope uses (explicitly not suitable for)	

3. Training and Input Data

Data sources	
Data time period / version	
Personal or special category data (Yes / No; detail)	
DPIA reference	

Data lineage / provenance location	Document data provenance and the controlled record. Where imported lineage is used, cite its verified AIG-INV-04 Evidence ID, source run or dataset and observed time.
------------------------------------	--

4. Performance

Key metrics and measured values	
Evaluation dataset	
Performance across relevant groups	

5. Fairness and Bias Testing

Protected characteristics tested	
Testing approach	
Results / disparities identified	
Mitigations applied	

6. Limitations and Risks

Known limitations	
Failure modes / edge cases	
Key residual risks	

7. Human Oversight

Human review / override mechanism	
Escalation route	

8. Agentic Capability and Action Scope (complete where the model can take actions)

Tools, systems and actions available to the model	List every tool, system or interface the model can invoke and the actions it can perform
Approved action scope	The actions the model is authorised to take and the systems on which it may act (Playbook §4.7.7)

Human control level per action capability	For each action capability: human executes / human approves each action / acts within defined bounds with monitoring / fully autonomous
Full agent composition	Held in the Agent Record (ASBOM): models, tools, interfaces, memory, sub-agents and the authority envelope. This card records the model; see the Agent, Interface and Dataset records.

9. Governance and Approval

Risk tier (Chapter 4) / governance priority	
DPIA / EIA references	
Delegated decision (decision-maker and date; AIG-DEC-03)	
Monitoring arrangements (metrics, thresholds, cadence)	

This card supports the transparency, explainability and accountability requirements of the Council AI Governance Playbook. An incomplete card indicates the model is not yet ready for governance approval.

10. Evidence and Assurance

Control field	Controlled value / evidence reference
Evidence pack location and version	Controlled repository/link, evidence-pack version and cut-off date. Where imported evidence is used, add its verified AIG-INV-04 Evidence ID and gate snapshot ID.
Performance and bias-test artefacts	Test report, dataset/version, results and acceptance thresholds. Where imported results are used, cite the verified Evidence ID and observed time.
Privacy, security and impact-assessment references	DPIA / EIA / HRA / RAIA / security review references, as applicable
Governance Decision Record	Decision record ID, approval conditions and date
Assurance reviewer and date	Name, role, review date and findings/conditions reference

11. Change, Rollback and Retirement

Control field	Controlled value / evidence reference
Current model, configuration, prompt and data versions	Record exact deployed versions and any available registry reference. Where imported evidence is used, cite its verified Evidence ID and observed snapshot. Review changes before approval.
Material-change triggers	Define changes that require reassessment, re-testing and governance re-approval
Model change history	Date, change, reason, test evidence, decision record and approver
Rollback control	Last approved version, rollback procedure, accountable owner and latest test date/result
Retirement / safe-withdrawal criteria	Triggers, service-continuity plan and authority to suspend or retire
Retirement and records	Retirement date; model/data disposal or return; retained evidence; Register status update

Document control

Administrative control record. This artefact remains a working draft until formally approved.

Field	Controlled value
Document ID	AIG-ASS-09
Version	1.5 draft
Status	Draft for Council review
Document owner	AI Governance Lead
Approving body	Pending — competent Council authority under confirmed Constitution / Scheme of Delegation; AI Assurance Board provides assurance and is not presumed to hold approval authority.
Approval date	Not yet approved
Effective date	Not yet effective
Next review date	To be set on approval: 12 months after the effective date, or sooner following material change
Classification	OFFICIAL

Change history

Version	Date	Change summary	Owner
1.0 draft	6 August 2026	Document-control baseline established following full-suite quality assurance; no substantive content change.	AI Governance Lead
1.1 draft	6 August 2026	QA revision: added evidence-pack, model-change, rollback and retirement controls.	AI Governance Lead

Version	Date	Change summary	Owner
1.2 draft	29 September 2026	QA remediation — applied findings QA-185, QA-200, QA-204, QA-221 (see QA change log). Final tidy-up: self-referencing cross-reference removed from the standards row(s) REQ-021, REQ-022.	AI Governance Lead
1.3 draft	30 September 2026	v3.9.1: standards cross-reference note updated to the 54-entry requirement inventory (AIG-AIMS-05 REQ-054, Civil Contingencies Act 2004, added).	AI Governance Lead
1.4 draft	1 October 2026	v3.9.3: standards cross-reference note cites the 55-entry requirement inventory (AIG-AIMS-05 REQ-055, NIST AI 600-1 Generative AI Profile, added).	AI Governance Lead
1.5 draft	2 October 2026	v3.9.5: standards cross-reference note cites the 56-entry requirement inventory (AIG-AIMS-05 REQ-056, SI 2026/425 ICO statutory code of practice on AI and automated decision-making, added). Proposed — for Council confirmation.	AI Governance Lead

NOT APPROVED — PROPOSED INTEGRATION ANNEX

Integration annex — Model-card evidence, version and lifecycle joins

DRAFT FOR COUNCIL REVIEW • 25 September 2026 • Not approved for operational use

Model card remains native

Retain one native card per model and record the exact model, configuration, prompt and relevant data versions; do not collapse multiple models into one undocumented entry. After registration, use the AIR-ID in AIG-INV-04. Index supporting evidence in the proposed AIG-INV-04 Evidence Index by Evidence ID, source/version, observed or verification date and location. Cite measured results only with their test artefact, dataset/version, conditions and limitations; this annex supplies none.

Change, monitoring and decision authority

For a changed model, supplier, configuration, prompt, data version, capability, or action scope, record the change and reassessment evidence in the card, update current-state/change records in AIG-INV-04, and link operational monitoring/review in AIG-OPS-02. An AIG-DEC-04 reference may identify only an actual dated Gate Event or Condition (ID and date). Approval and conditions are recorded in AIG-DEC-03 or authoritative forum minutes; AIG-DEC-04 is not an approval.

Agent capability and affected people

For agentic capability, the AIG-AGT-04 ASBOM is the permissions/delegation record: link the exact AG-ID and version. AIG-AGT-05 is a derived graph, not a separate source of authority; AIG-AGT-06 is the action-specific decision record. Link AIG-OPS-04 for contestability where model outputs make or materially influence individual decisions. These references do not confer permission, establish an outcome, or demonstrate effective controls.

Draft status

No performance finding, validation result, legal applicability, operational state, approval or publication is established by this annex. Keep unsupported fields explicit as unknown/not evidenced rather than infer completion.

Standards and legal applicability — proposed QA cross-reference

PROPOSED QA INDEX — NOT APPROVED LEGAL ADVICE OR EVIDENCE OF COMPLIANCE. The listed instruments may be relevant to this artefact, but applicability depends on the case, service and Council delegation. Legislation, official guidance, Council policy and voluntary standards are different authority classes. The full 56-entry requirement inventory, authoritative URLs, validation owners and current versions are in AIG-AIMS-05; source checks are in AIG-AIMS-13. Do not treat a template, plan or cross-reference as completed assessment, adopted policy, ISO conformity or legal sign-off.

Requirement / instrument	Why check it here	Source
REQ-001 — UK General Data Protection Regulation (UK GDPR), as amended	Where personal data is processed; test each UK GDPR provision separately.	https://www.legislation.gov.uk/eur/2016/679/contents
REQ-021 — Copyright, Designs and Patents Act 1988	Screen the legislation requirement against this artefact's actual use and record evidence/version.	https://www.legislation.gov.uk/ukpga/1988/48/contents
REQ-022 — Copyright and Rights in Databases Regulations 1997	Screen the legislation requirement against this artefact's actual use and record evidence/version.	https://www.legislation.gov.uk/uksi/1997/3032/contents

REQ-041 — ISO/IEC 23894:2023 AI — Guidance on risk management	Voluntary risk-management guidance; proposed alignment, not statutory conformity.	https://www.iso.org/standard/ 77304.html
---	---	--